

# Oracle Developer Essentials: Tables and Indexes

Introduction

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Course Outline

Defining  
tables

Table  
constraints

Storage  
options

Managing  
existing  
tables

Special table  
types

Indexes

Permissions

Good  
practices

# Oracle Essentials Course Sequence

The diagram consists of three rounded rectangular boxes arranged horizontally. A light blue arrow points from the first box to the second, and another light blue arrow points from the second box to the third. The first box is green and contains the text 'Tables and Indexes'. The second box is teal and contains the text 'Data Types'. The third box is magenta and contains the text 'Views, Synonyms and Triggers'.

Tables and Indexes

Data Types

Views,  
Synonyms  
and Triggers

# Database Lifetime and Usage

## Used by **multiple** applications

---

- Most data has multiple **consumers**
- **Many applications** will be connected over the **lifetime** of the database

## Used for **long periods** of time

---

- Often used for **five** or more years
- **Historical** data is often important

# Motivation

Understand all the **available options**

Build databases that are **easy to use**

Build databases that **perform well**

# *Terminology*



# Database Objects

*Generic term for a structure owned by a user in an Oracle database*

Tables

Views

Indexes

Synonyms

Sequences

Types

Functions

Procedures

Triggers

# Users in Oracle

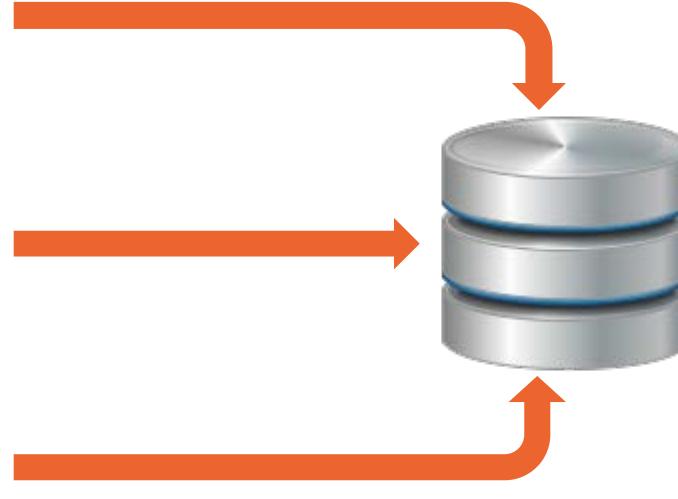
User

A user within the Oracle database

dberry  
(individual user)

student  
(application user)

student\_web  
(application user)



# Schemas

Definition

All of the objects that are owned by a user

As a Container

Often correlate to an application or business function

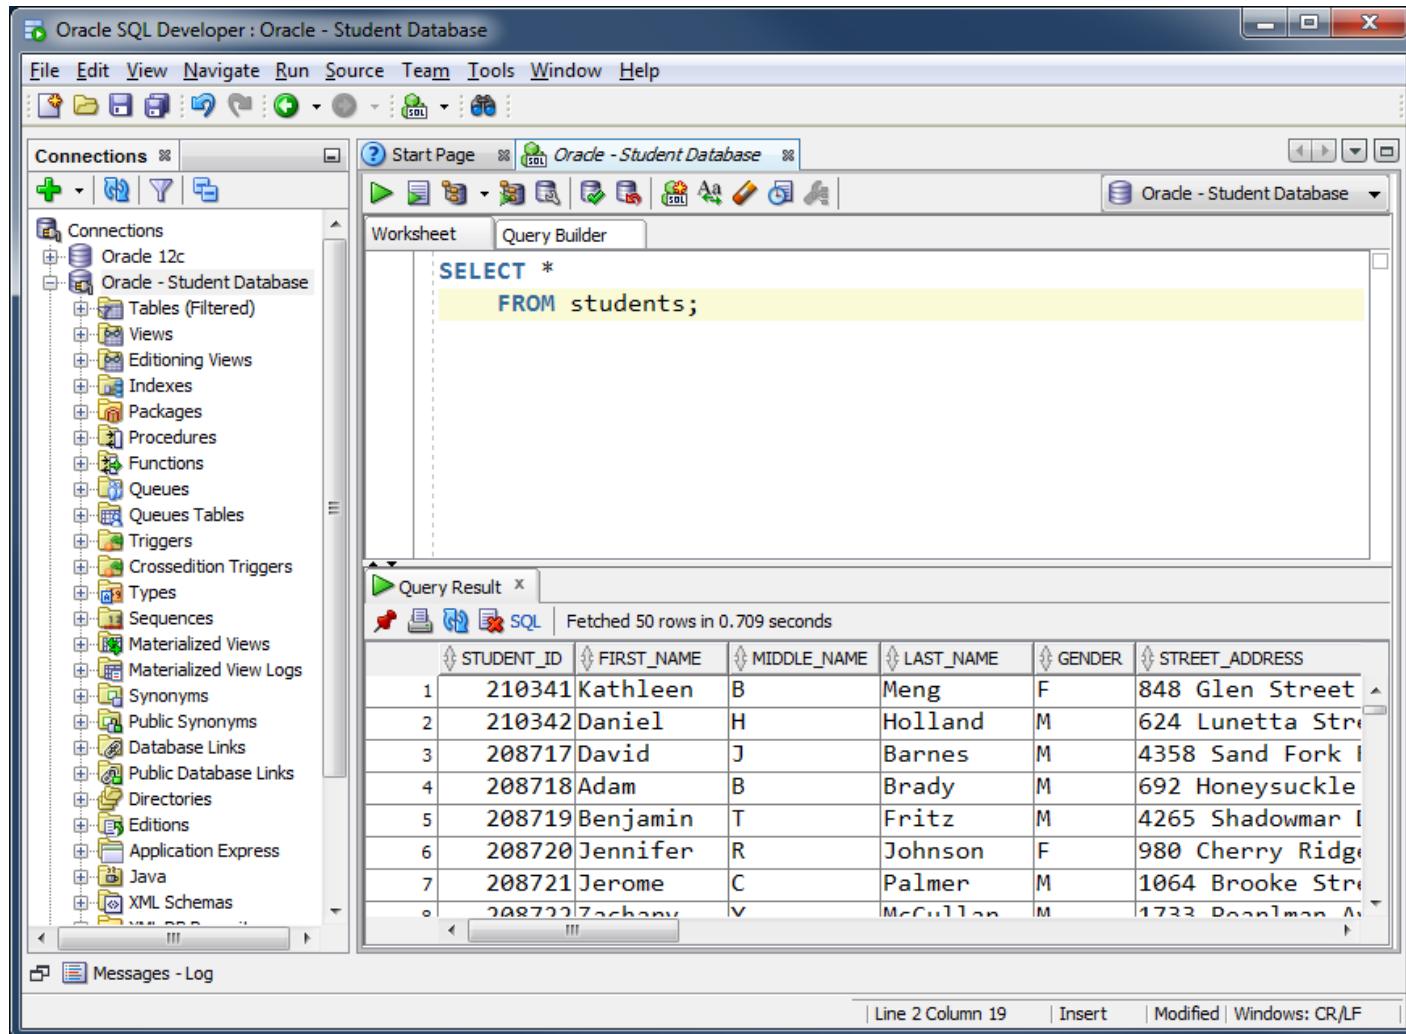
# Accessing Objects in a Different Schema

```
SELECT  
    department_code, course_number, name, description  
FROM student.courses
```



*Schema name must be included  
when accessing the table as a  
different user*

# Oracle SQL Developer

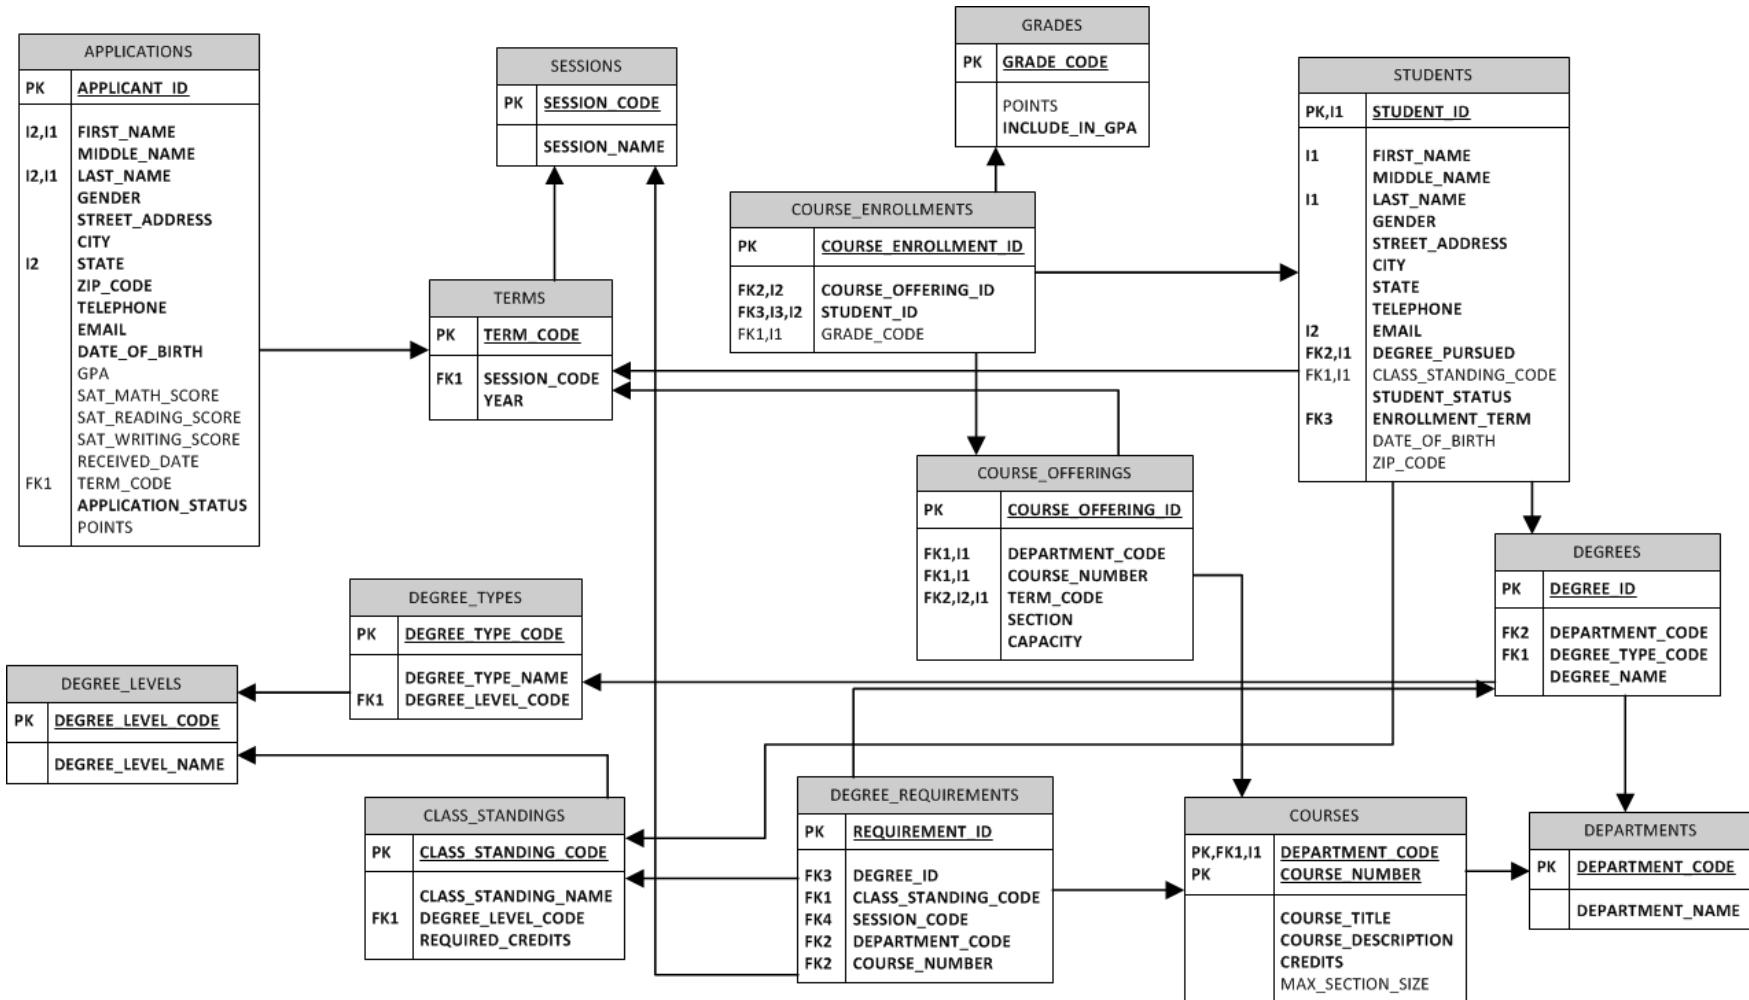


<http://bit.ly/OracleSqlDeveloper>

# Course Conventions

- ◆ Course will use SQL statements to demonstrate capabilities
- ◆ SQL syntax better communicates different options
- ◆ Can easily transition to graphical tools later
- ◆ Syntax of commands important for reading scripts

# Course Resources



<http://bit.ly/OracleCourseSchema>

# Oracle Resources

Oracle 12c Documentation Home

<http://bit.ly/Ora12cDocs>

Oracle Developer Virtual Machines

<http://bit.ly/OracleVms>

# Summary

**Course overview and motivation**

**Common Oracle terms**

**Oracle resources**

# Table Basics

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Introduction to Tables

student_id	first_name	last_name	phone	email	major_id
208718	Adam	Brady	360-592-7886	AdamBBrady@dayrep.com	3
208723	George	Ward	216-952-0212	GKWard@superrito.com	5
208728	Marcia	Garcia	630-239-4325	MarciaDGarcia@rhyta.com	8
208741	Walter	Stumpf	419-500-5489	WalterCStumpf@einrot.com	12
208749	Brenda	Brown	979-859-7105	BrendaBrown@dayrep.com	5
208755	Christopher	Glenn	803-980-0813	ChrisGlenn@superrito.com	5
209628	Lisa	Shifflett	229-409-0015	lshifflett@jourrapide.com	4

# Tables Roadmap

## Table Basics (this module)

---

Create table syntax  
Naming rules  
Null values  
Default values

## Table Constraints (next module)

---

Primary keys  
Foreign keys  
Check constraints

## Table Storage (in two modules)

---

Storage characteristics  
Storage options

# Naming Rules

## Maximum Length

30 characters for all objects

## Uniqueness

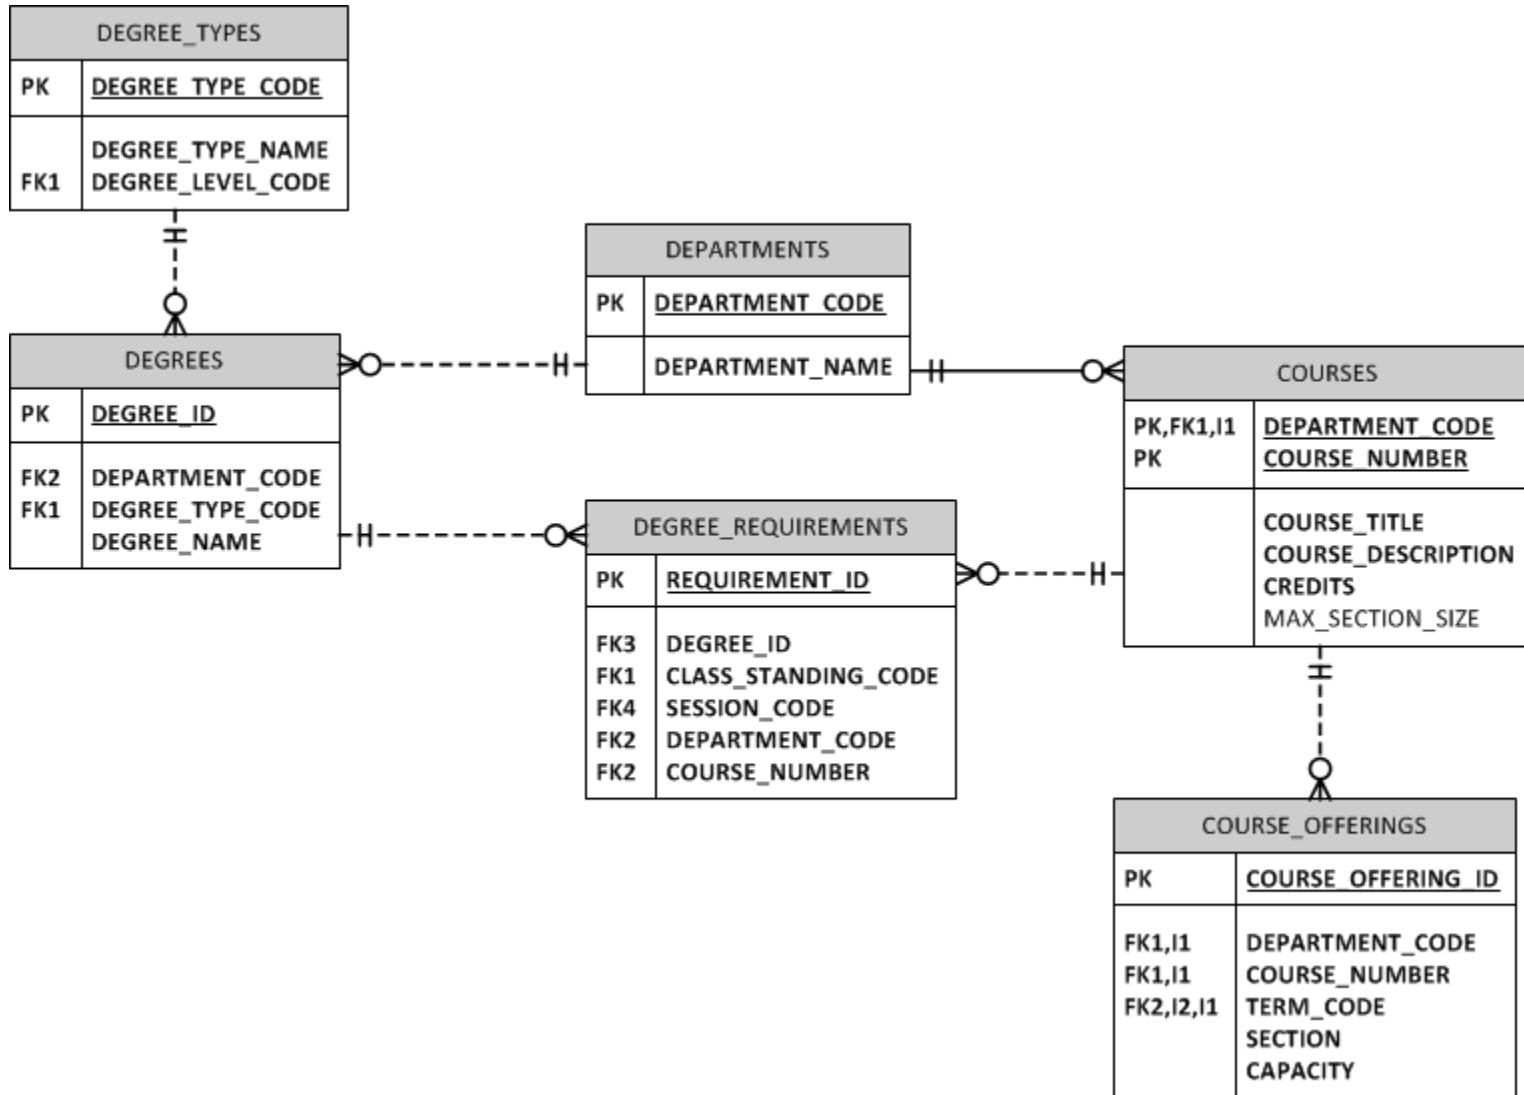
Tables, views must be unique within schema

Columns must be unique within a table

## SQL Reserved Words

Possible to use but discouraged

# Tables Represent Business Data



# Basic Create Table Syntax

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses PRIMARY KEY
        (department_code, course_number),
    CONSTRAINT fk_courses_department_code FOREIGN KEY
        (department_code) REFERENCES departments (department_code)
    CONSTRAINT ck_courses_course_number
        CHECK (course_number BETWEEN 100 AND 999)
)
TABLESPACE users
PCTFREE 75;
```

# **Naming Tables and Columns**

Good names **communicate intent**

Poor names **lead to confusion**

# Naming Rules

## Maximum Length

30 characters for all objects

## Uniqueness

Tables, views must be unique within schema

Columns must be unique within a table

## SQL Reserved Words

Possible to use but discouraged

# Naming Rules

## Maximum Length

- 30 characters for all object types

## Uniqueness

- Tables, views, indexes must be unique within a user schema
- Columns must be unique within a table or view

## SQL Reserved Words

- Possible to use but discouraged

# Unquoted Names

- Most common way of naming objects

```
CREATE TABLE zip_codes
(
    zip_code      VARCHAR2(5)      NOT NULL,
    city          VARCHAR2(30)     NOT NULL,
    state         VARCHAR2(30)     NOT NULL
);
```

- Allowable characters are

- Alphanumeric characters
- Underscore (\_) character
- \$ and # characters

- Names will be case insensitive

# Unquoted Names Are Case Insensitive

```
CREATE TABLE zip_codes
(
    zip_code      VARCHAR2(5)      NOT NULL,
    city          VARCHAR2(30)     NOT NULL,
    state         VARCHAR2(30)     NOT NULL
);
```

## Allowed SQL Statements

```
-- This is OK
select ZIP_CODE, CITY, STATE
    from ZIP_CODES;

-- So is this
SELECT Zip_Code, City, State
    FROM Zip_Codes;
```

# Quoted Identifiers

- Allows a wider range of options

```
CREATE TABLE “ZipCodes”
(
    “zip code”          VARCHAR2(5)      NOT NULL,
    “city.name”         VARCHAR2(30)     NOT NULL,
    “state-abbr”        VARCHAR2(30)     NOT NULL
);
```

- You must also use quoted names in SQL statements

```
SELECT
    “zip code”, “city.name”, “state-abbr”
FROM “ZipCodes”
WHERE “zip code” = ‘54911’
```

# Defining Columns

```
CREATE TABLE students
(
    student_id      NUMBER(6)      NOT NULL,
    first_name      VARCHAR2(30)    NOT NULL,
    last_name       VARCHAR2(30)    NOT NULL,
    email_address   VARCHAR2(128)   NOT NULL,
    phone           VARCHAR2(20)    NULL,
    city            VARCHAR2(30)    NULL,
    state           VARCHAR2(2)     NULL,
    status_code     VARCHAR2(1)     DEFAULT 'A' NOT NULL,
    CONSTRAINT pk_students PRIMARY KEY (student_id)
);
```

# Column Definition Rules

## Number of Columns per Table

- Maximum of 1000 columns
- Over 255 columns results in row chaining

## Column Naming Rules

- 30 characters maximum
- Names must be unique within a table
- Names can be reused in different tables

## Minimum Information

- Name and datatype must be specified
- NULL and default clauses are optional

# NULL Column Values

```
CREATE TABLE students
(
    student_id      NUMBER(6,0)      NOT NULL,
    first_name      VARCHAR2(30)     NOT NULL,
    last_name       VARCHAR2(30)     NOT NULL,
    city            VARCHAR2(30)      NULL,
    state           VARCHAR2(2)       NULL,
    date_of_birth   DATE
);
```

- In character based columns (CHAR, VARCHAR, NCHAR, NVARCHAR)
  - Empty string is treated as a NULL value

# Default Values

```
CREATE TABLE students
(
    student_id          NUMBER(6,0)      NOT NULL,
    first_name          VARCHAR2(30)     NOT NULL,
    last_name           VARCHAR2(30)     NOT NULL,
    city                VARCHAR2(30)      NULL,
    state               VARCHAR2(2)       NULL,
    date_of_birth       DATE,
    enrollment_status   VARCHAR2(1)      DEFAULT 'E' NOT NULL
);
```

# **NULL, NOT NULL or DEFAULT VALUE**

---

That is the question

# Data You Do Not Have

*Using NULL values*

user_id	email_address	first_name	last_name	phone_number	city
8385930	AdamBBrady@dayrep.com	Adam	Brady	303-555-1212	Denver
8385931	GKWard@superrito.com	George	Ward	<NULL>	Phoenix
8385932	MarciaDGarcia@rhyta.com	Marcia	Garcia	<NULL>	<NULL>
8385933	WalterCStumpf@einrot.com	Walter	Stumpf	414-555-8965	<NULL>

*Using default values*

user_id	email_address	first_name	last_name	phone_number	city
8385930	AdamBBrady@dayrep.com	Adam	Brady	303-555-1212	Denver
8385931	GKWard@superrito.com	George	Ward	000-000-0000	Phoenix
8385932	MarciaDGarcia@rhyta.com	Marcia	Garcia	000-000-0000	UNKNOWN
8385933	WalterCStumpf@einrot.com	Walter	Stumpf	414-555-8965	UNKNOWN

# An Event That Has Not Occurred

*Using NULL values*

order_id	customer_id	order_status_code	order_date	ship_date
567893	832	S	2014-08-03	2014-08-05
56784	944	S	2014-08-03	2014-08-07
567895	968	P	2014-08-05	<NULL>
567896	1024	N	2014-08-06	<NULL>

*Using default values*

order_id	customer_id	order_status_code	order_date	ship_date
567893	832	S	2014-08-03	2014-08-05
56784	944	S	2014-08-03	2014-08-07
567895	968	P	2014-08-05	1900-01-01
567896	1024	N	2014-08-06	1900-01-01

# Use Defaults To Indicate an Initial State

*Using default values*

member_id	email_address	first_name	last_name	phone_number	members hip_level
8385930	AdamBBrady@dayrep.com	Adam	Brady	303-555-1212	P
8385931	GKWard@superrito.com	George	Ward	602-555-3547	F
8385932	MarciaDGarcia@rhyta.com	Marcia	Garcia	608-555-2575	F
8385933	WalterCStumpf@einrot.com	Walter	Stumpf	414-555-8965	P

*Using NULL values*

member_id	email_address	first_name	last_name	phone_number	members hip_level
8385930	AdamBBrady@dayrep.com	Adam	Brady	303-555-1212	P
8385931	GKWard@superrito.com	George	Ward	602-555-3547	
8385932	MarciaDGarcia@rhyta.com	Marcia	Garcia	608-555-2575	
8385933	WalterCStumpf@einrot.com	Walter	Stumpf	414-555-8965	P

# NULL versus Default Values

## Know your data

- What does the data mean
- How is it used

## What communicates intent most clearly?

- Would another user draw the same conclusion?

# Virtual Columns Introduction

## Normal column

---

Data is **physically** stored  
on disk by Oracle

## Virtual column

---

Value is **computed** from  
other columns in table

# Virtual Column Use Case

```
CREATE TABLE orders
(
    order_id      NUMBER(10)      NOT NULL,
    order_date    DATE           NOT NULL,
    customer_id   NUMBER(10)      NOT NULL,
    subtotal      NUMBER(10,2),
    tax           NUMBER(10,2),
    shipping      NUMBER(10,2),
    grand_total   NUMBER(10,2)
);
```

# Defining a Virtual Column

```
CREATE TABLE orders
(
    order_id      NUMBER(10)      NOT NULL,
    order_date    DATE           NOT NULL,
    customer_id   NUMBER(10)      NOT NULL,
    subtotal      NUMBER(10,2),
    tax           NUMBER(10,2),
    shipping      NUMBER(10,2),
    grand_total   NUMBER(10,2)
                           AS (subtotal + tax + shipping) VIRTUAL
);
```

# Defining a Virtual Column With a Function

```
CREATE TABLE students5
(
    student_id      NUMBER(10)      NOT NULL,
    first_name      VARCHAR2(30)     NOT NULL,
    middle_name     VARCHAR2(30)     NOT NULL,
    last_name       VARCHAR2(30)     NOT NULL,
    email_address   VARCHAR2(60)     NOT NULL,
    email_domain   VARCHAR2(60) AS (
        SUBSTR(email_address, INSTR(email_address, '@', 1, 1)+1)
    ) VIRTUAL
);
```

# **Virtual Column Properties**

- ◆ **Value is computed in result set of query**
- ◆ **Cannot INSERT into or UPDATE virtual columns**
- ◆ **Can only make use of columns in the same table**
- ◆ **Indexes can be created over virtual column values**

# **Virtual Columns Summary**

**Useful for situations where a derived value is needed**

# **Summary**

**CREATE TABLE**  
Syntax

Naming Rules

**NULLS** and  
Default Values

# Table Constraints

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# **Table Constraints Overview**

- **Table constraints**
  - Define rules your data must follow
- **Examples**
  - Every student must have a unique student id number
  - Every course must be associated with a department

# **Module Outline**

Primary Keys

Foreign Keys

Check Constraints

# Primary Keys

- Uniquely identifies each row in a table
- Can be defined over one or more columns
- Cannot contain a NULL value in any primary key column

The diagram illustrates a primary key constraint. A solid line connects the primary key column 'STUDENT\_ID' in the STUDENTS table to the foreign key column 'STUDENT\_ID' in the COURSE\_ENROLLMENTS table. Below the tables, there is a symbol consisting of a vertical dashed line with a circle at the bottom, indicating a relationship between the two tables.

STUDENTS	
PK,I1	<u>STUDENT_ID</u>
I1	FIRST_NAME
I1	MIDDLE_NAME
I1	LAST_NAME
I1	GENDER
I1	STREET_ADDRESS
I1	CITY
I1	STATE
I1	TELEPHONE
I2	EMAIL
FK2,I1	DEGREE_PURSUED
FK1,I1	CLASS_STANDING_CODE
FK1,I1	STUDENT_STATUS
FK3	ENROLLMENT_TERM
FK3	DATE_OF_BIRTH
FK3	ZIP_CODE

COURSE_ENROLLMENTS	
PK	<u>COURSE_ENROLLMENT_ID</u>
FK2,I2	COURSE_OFFERING_ID
FK3,I3,I2	STUDENT_ID
FK1,I1	GRADE_CODE

# Primary Key Types

## Natural Key

Combination of naturally occurring columns that form a unique key

## Surrogate Key

A unique value is generated for each row

***The real reason I favor surrogate keys....***



***I am a terrible typist!***

# Primary Key Definition

```
CREATE TABLE departments
(
    department_code      VARCHAR2(2)      NOT NULL PRIMARY KEY,
    department_name     VARCHAR2(64)      NOT NULL
);
```

# Primary Key Using Constraint Clause

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses
        PRIMARY KEY (department_code, course_number)
);
```

# Add a Primary Key to an Existing Table

```
ALTER TABLE courses
    CONSTRAINT pk_courses
    ADD PRIMARY KEY (department_code, course_number);
```

# Primary Key Recommendations

Make sure **every table**  
has a **primary key**  
defined

Values of a **primary**  
**key** should be  
**immutable**

# Oracle ROWID

- Represents the address of the row in the table

```
SQL> SELECT rowid, zip_code, city, state
  2      FROM zip_codes;
```

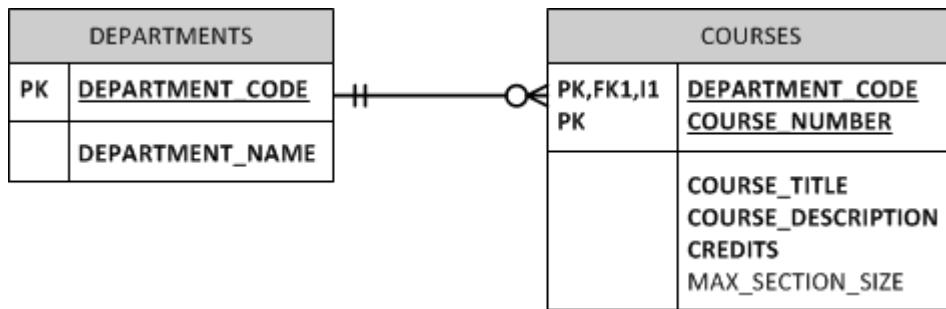
ROWID	ZIP_C CITY	ST
AAA...zeAAA	11201 Brooklyn	NY
AAA...zeAAB	75201 Dallas	TX
AAA...zeAAC	80401 Golden	CO
AAA...zeAAD	92101 San Diego	CA

# Foreign Keys

---

Put the relational in relational databases

# What Is a Foreign Key



For a record to exist in the **COURSES** table, it must contain a valid **DEPARTMENT\_CODE** value from the **DEPARTMENTS** table

# Defining a Foreign Key

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses
        PRIMARY KEY (department_code, course_number),
    CONSTRAINT fk_courses_department_code
        FOREIGN KEY (department_code)
        REFERENCES departments (department_code),
);
;
```

# Add a Foreign Key to an Existing Table

```
ALTER TABLE courses
  CONSTRAINT fk_courses_department_code
  ADD FOREIGN KEY (department_code)
  REFERENCES departments (department_code);
```

# On Delete Cascade

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses
        PRIMARY KEY (department_code, course_number),
    CONSTRAINT fk_courses_department_code
        FOREIGN KEY (department_code)
            REFERENCES departments (department_code)
            ON DELETE CASCADE
);
```

# On Delete Cascade

*departments*

department_code	name
MA	Math
PH	Physics
CS	Computer Science



*courses*

department_code	course_number	Title
MA	101	Calculus 1
MA	102	Calculus 2
PH	101	Physics 1
PH	102	Physics 2
CS	101	Intro to Programming



```
DELETE FROM departments  
WHERE department_code = 'PH';
```

# On Delete Cascade

- Without on delete cascade

```
DELETE FROM courses WHERE department_code = 'PH';
DELETE FROM departments WHERE department_code = 'PH';
COMMIT;
```

- To delete a primary key value, no child records can exist

# Deferred Constraints

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses
        PRIMARY KEY (department_code, course_number),
    CONSTRAINT fk_courses_department_code
        FOREIGN KEY (department_code)
        REFERENCES departments (department_code)
        DEFERRABLE
        INITIALLY IMMEDIATE
        -- INITIALLY DEFERRABLE
);
```

# Using a Deferred Constraint

*departments*

department_code	name
MA	Math
PH	Physics
CS	Computer Science

*courses*

department_code	course_number	Title
MA	101	Calculus 1
MA	102	Calculus 2
PH	101	Physics 1
PH	102	Physics 2
CS	101	Intro to Programming

```
set constraint
fk_courses_department_code deferred;

UPDATE departments
SET department_code = 'CO'
WHERE department_code = 'CS';

UPDATE courses
SET department_code = 'CO'
WHERE department_code = 'CS';
COMMIT;
```

# **Foreign Key Options Summary**

- **Why are you having to cascade deletes or defer constraints?**
  - One time cleanup or typical business operation
- **Using on delete cascade**
  - Is deleting data really the correct decision?
  - Can the parent record be marked 'inactive'
- **Using deferred constraints**
  - Consider a surrogate key for a primary key

# Check Constraints

Check that a value is  
within a given **range**

**Compare** two values  
in **different columns**

Validate the format of  
a value using a  
**regular expression**

# Check If Value is Within a Range

```
CREATE TABLE courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses
        PRIMARY KEY (department_code, course_number),
    CONSTRAINT fk_courses_department_code
        FOREIGN KEY (department_code)
        REFERENCES departments (department_code),
    CONSTRAINT ck_courses_course_number
        CHECK (course_number BETWEEN 100 AND 999)
);
```

# Check If Value is Within a Domain

```
CREATE TABLE states
(
    state_code          VARCHAR2(2)      NOT NULL,
    state_name         VARCHAR2(30)      NOT NULL,
    region             VARCHAR2(2)      NOT NULL,
    CONSTRAINT pk_states
        PRIMARY KEY (state_code),
    CONSTRAINT ck_state_regions
        CHECK (region IN ('NE', 'SE', 'MW', 'SC', 'NW', 'SW'))
);
```

# Emulate a Boolean Column

```
CREATE TABLE students
(
    student_id      NUMBER(6)      NOT NULL,
    first_name      VARCHAR2(30)    NOT NULL,
    last_name       VARCHAR2(30)    NOT NULL,
    email_address   VARCHAR2(128)   NOT NULL,
    likes_ice_cream NUMBER(1)      NULL,
    CONSTRAINT pk_students PRIMARY KEY (student_id),
    CONSTRAINT ck_studens_ice_cream
        CHECK (likes_ice_cream IN (0,1))
);
```

# Compare Two Values

```
CREATE TABLE Orders
(
    order_id      NUMBER(9)      NOT NULL,
    customer_id   NUMBER(6)      NOT NULL,
    order_date    DATE          NOT NULL,
    ship_date     DATE          NULL,
    order_status  VARCHAR2(1)    NOT NULL,
    CONSTRAINT pk_orders
        PRIMARY KEY (order_id),
    CONSTRAINT ck_orders_order_ship_date
        CHECK (ship_date > order_date)
);
```

# Validate Format Using a Regular Expression

```
CREATE TABLE zip_codes
(
    zip_code      VARCHAR2(5)      NOT NULL,
    city          VARCHAR2(30)     NOT NULL,
    state_code    VARCHAR2(30)     NOT NULL,
    CONSTRAINT pk_codes
        PRIMARY KEY (state_code),
    CONSTRAINT ck_zip_code_format
        CHECK (REGEXP_LIKE (zip_code, '^[0-9]{5}$') )
);
```

# Add a Constraint to an Existing Table

```
ALTER TABLE zip_codes
    ADD CONSTRAINT ck_zip_codes_format
    CHECK (REGEXP_LIKE (zip_code, '^[0-9]{5}$') );
```

# Disabling and Enabling Constraints

## Disabling a constraint

```
ALTER TABLE courses  
    DISABLE CONSTRAINT fk_courses_department_code;
```

## Enabling a constraint

```
ALTER TABLE courses  
    ENABLE CONSTRAINT fk_courses_department_code;
```

# Constraints and Data Integrity

**Multiple options** are available for enforcing **data integrity**

**Data integrity** is critical to your **business data**

# Why Use Database Constraints

- Clean Data**
  - Required for many business functions
  - Can reduce the amount of coding needed to deal with “bad data”
- Consistent Enforcement of Constraints**
  - All applications must comply with database constraints

# Enforcing Constraints in Your Application

App is closer to user

- Faster feedback on invalid data
- Less frustrating to user

General good practice

- Mitigate security concerns
- Reduce crashes due to invalid data

# Constraints Strategy

## Database Constraints

- One set of tools in an overall strategy
- Make use of all tools available

## Validate both in Database and Applications

- Applications: provide immediate user feedback
- Database: consistently enforce rules across all applications

## Performance Impacts

- Constraints can be checked in milliseconds
- Bad data takes a long time to clean up

# **Summary**

Primary Keys

Foreign Keys

Check  
Constraints

# Table Storage Options

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Introduction

*students table*

student_id	first_name	last_name	phone
208718	Adam	Brady	360-592-7886
208723	George	Ward	216-952-0212
208728	Marcia	Garcia	630-239-4325
208741	Walter	Stumpf	419-500-5489
208749	Brenda	Brown	979-859-7105



# Module Outline

Database  
Blocks

How Oracle  
Reads Data

Tablespaces

High Water  
Mark

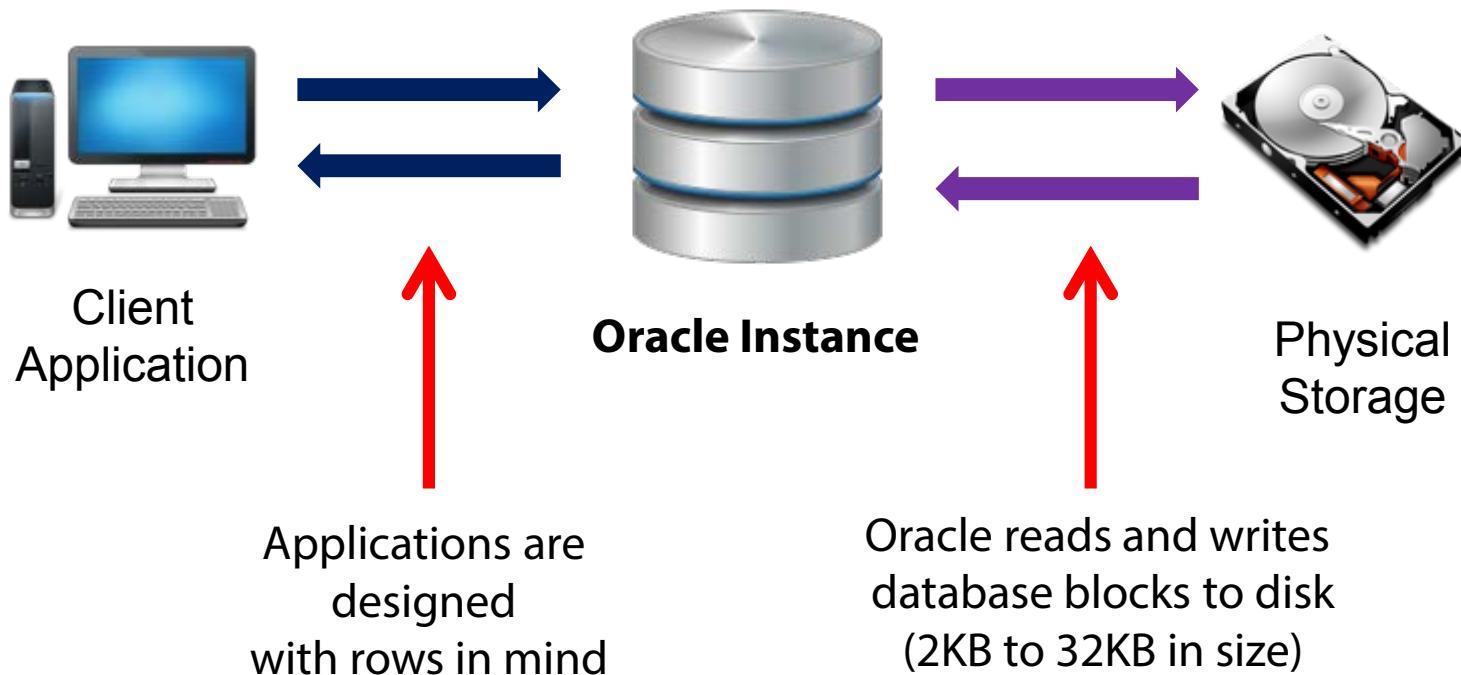
PCTFREE and  
PCTUSED

Row Migration

Freelists

Row Chaining

# Database Blocks



*A data block is like a shipping container*



# What Is In a Block

## Block Header

- Block type information
- Object assignment
- Row directory

## Data

- Actual data for your table

## Free Space

- Allows for updates to rows

## Database Block

### Block Header

Row

Row

Row

Row

Row

Free Space

Row

Row

Row

Free Space

# Oracle Storage Terminology

## Data Block

Smallest level of storage in Oracle

Can be between 2-32KB (8KB common)

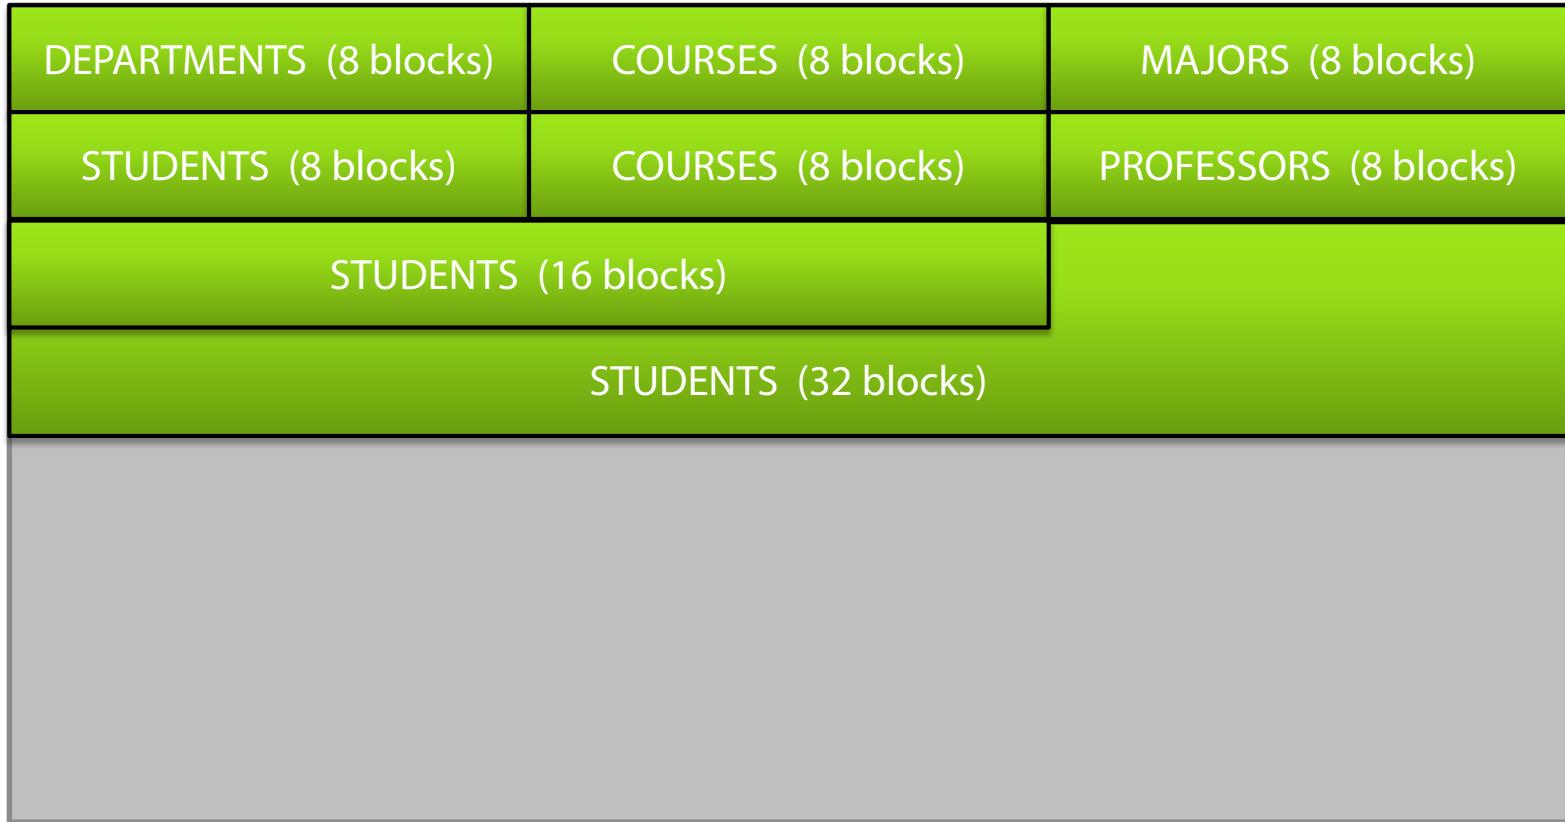
## Extent

Contiguous allocation of blocks for storing a specific type of data

## Segment

Set of extents allocated for a specific database structure

# Storage Allocation



# How Oracle Performs Data Access

**Full table scan**

**Index operation with  
row lookup**

# Full Table Scan

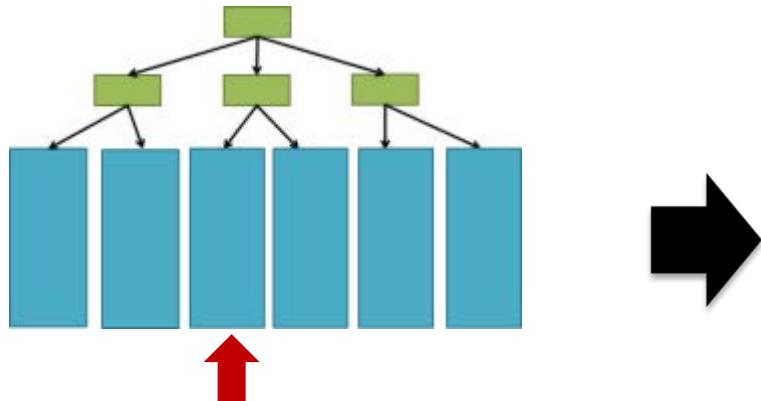
Table



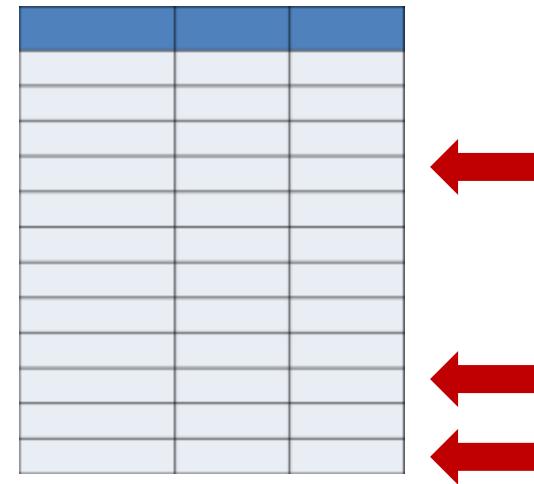
- Oracle scans through every **block** in the table
- Comparable to a linear search of an array

# Index Operation and Row Lookup

Index operation



Row lookup in table



*Oracle traverses the index tree to find matching keys*

*Oracle looks up the rows in the table by ROWID from index*

## Why Does This Matter?

Some situations **affect one data access** path,  
but **not the other**

# Table Data and Sort Order

Data in a table has no implied order

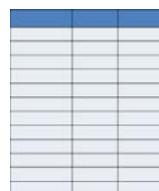
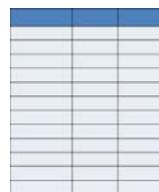
- Oracle will insert rows wherever they fit
- Physical order on disk != Order data was inserted in

Data returned from queries is not sorted

- Order of results is not order data was inserted
- Two different executions can return same results in different order

# Tablespace Introduction

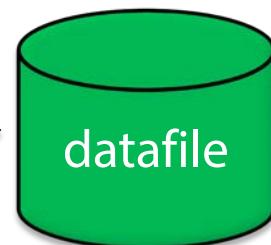
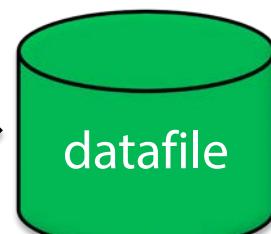
*Tables and indexes*



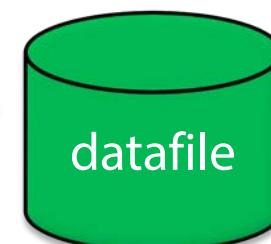
*Tablespace*

*Tablespace*

*Data files*



*datafile*



# Tablespace Block Size

Between 2 KB and 32 KB  
(8 KB is a typical size)

Useful for estimating number of rows per block

Need to know to avoid row chaining

# **Segment Space Management**

## **Manual Segment Space Management (MSSM)**

- Original space management implementation
- You are responsible for managing all parameters for space management

## **Automatic Segment Space Management (ASSM)**

- Introduced in Oracle 9i
- You specify value for PCTFREE
- Oracle manages all other parameters

# Oracle Contains Multiple Tablespaces

SYSTEM and  
SYSAUX

- Data about your Oracle instance
- Reserved for use by Oracle

TEMP

- Data in temporary tables
- Sorts and joins that do not fit in memory

User Tablespaces

- Data for your applications
- May be one per application, group of applications

# Default Tablespace

Every user has a  
**default** tablespace

Some objects may be  
created in an  
**alternate** tablespace

# Specifying a Tablespace for a Table

```
CREATE TABLE course_enrollments
(
    course_enrollment_id      NUMBER      NOT NULL,
    course_offering_id        NUMBER(20)   NOT NULL,
    student_id                 NUMBER(10)   NOT NULL,
    grade_code                  VARCHAR2(1)  NULL,
    CONSTRAINT pk_course_enrollments
        PRIMARY KEY (course_enrollment_id),
    CONSTRAINT fk_enrollments_offerings
        FOREIGN KEY (course_offering_id)
            REFERENCES course_offerings (course_offering_id),
    CONSTRAINT fk_enrollments_student
        FOREIGN KEY (student_id) REFERENCES students (student_id),
    CONSTRAINT fk_enrollments_grades
        FOREIGN KEY (grade_code) REFERENCES grades (grade_code)
)
TABLESPACE student_data;
```

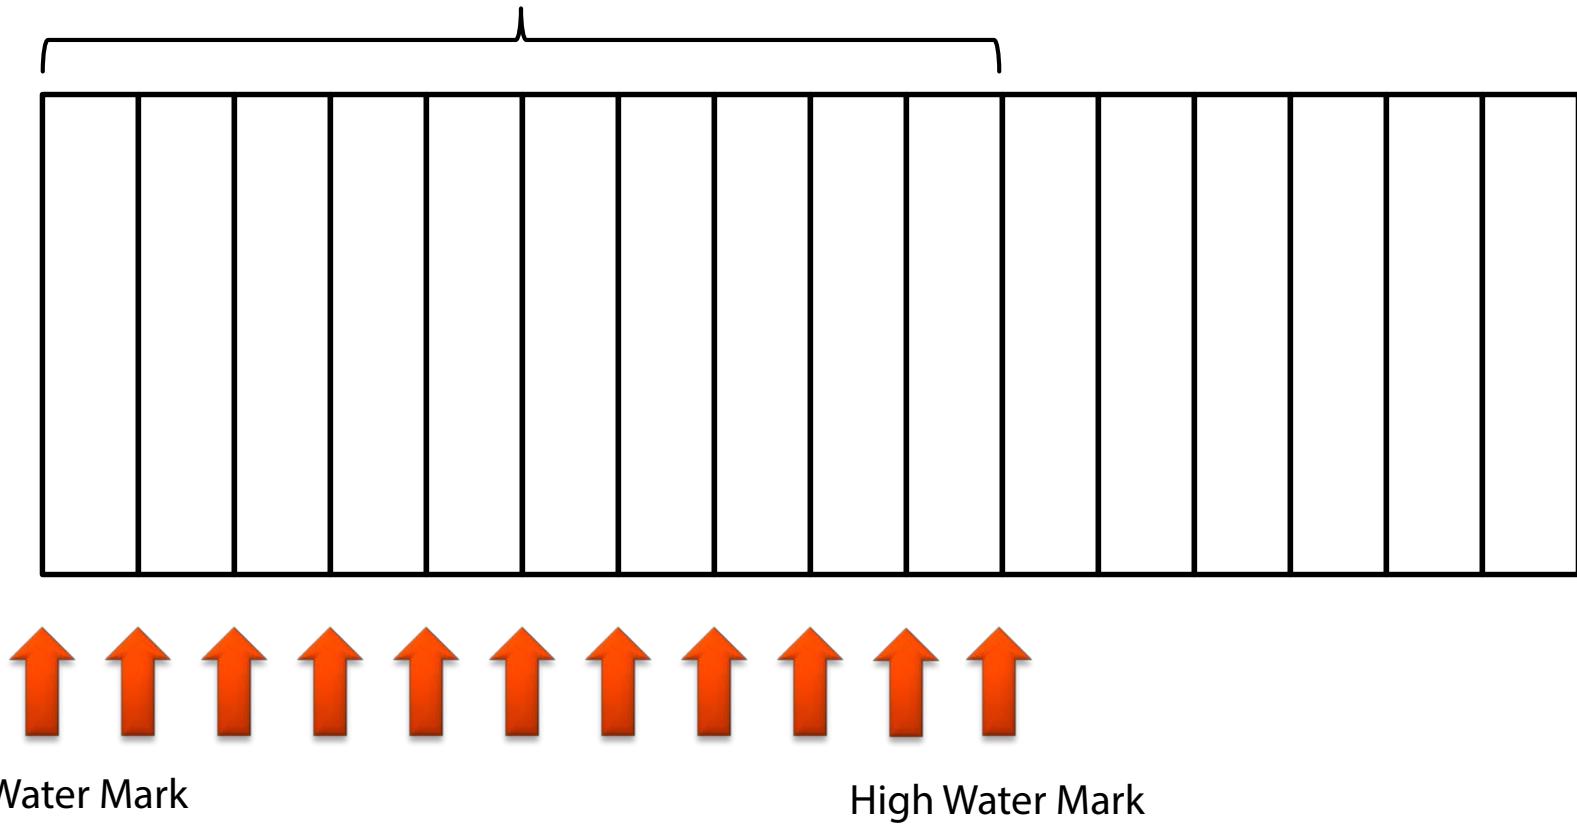
# **High Water Mark**

---

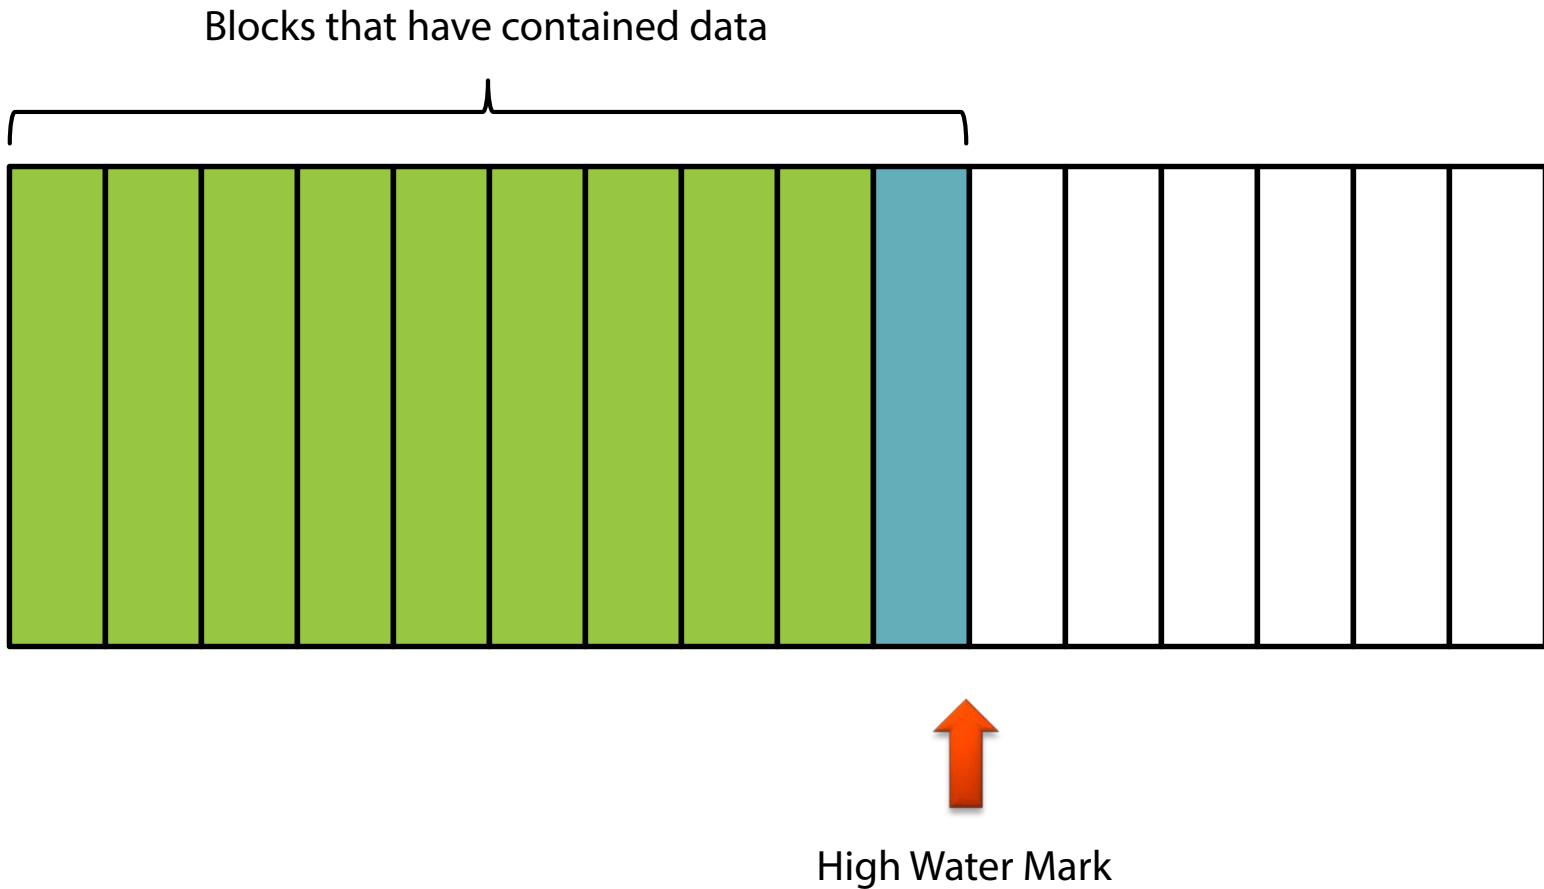
**Highest numbered block that has  
ever contained data in a table**

# High Water Mark

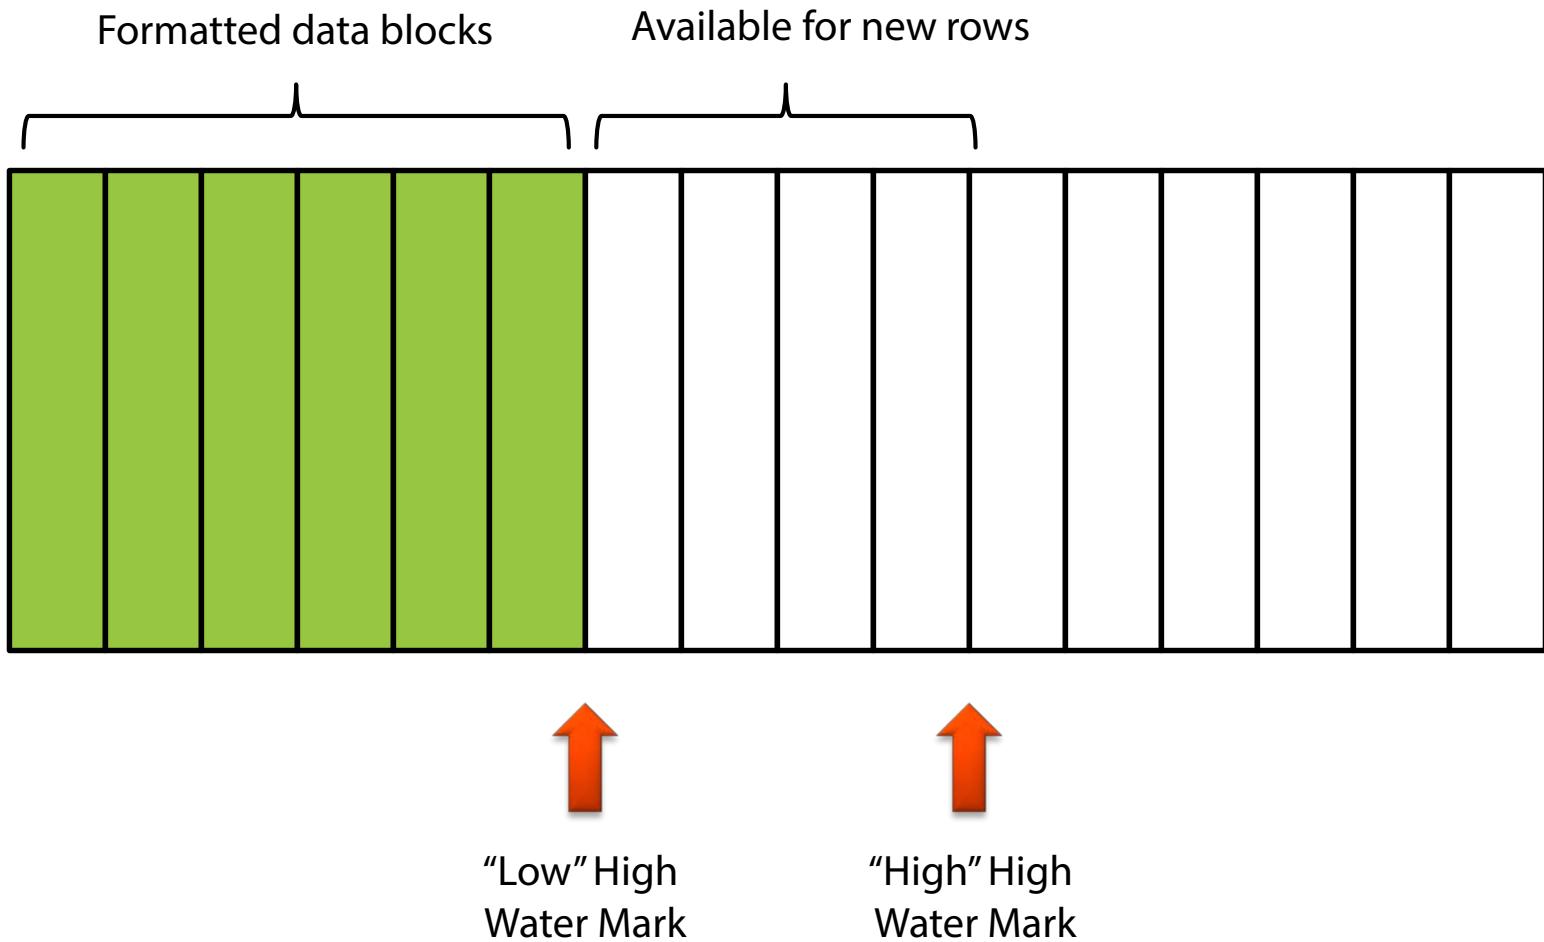
Blocks that have contained data



# High Water Mark – Deleting Rows



# High Water Mark – ASSM



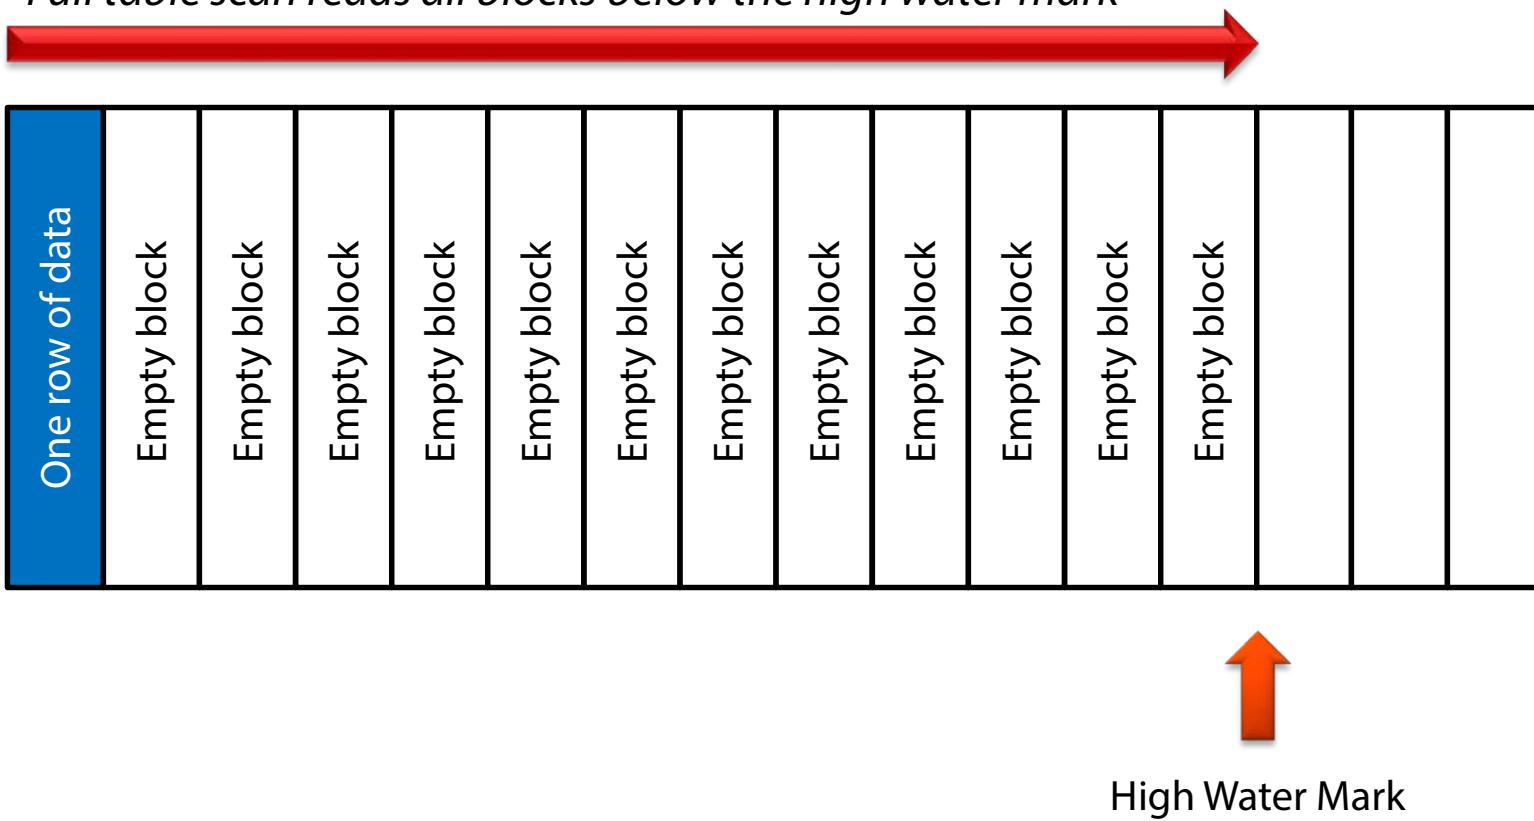
# High Water Mark Impacts

During a **full table scan**,  
Oracle reads all blocks  
below the **high water  
mark**

If many of these blocks  
are **empty**, this is a lot of  
**wasted work**

# High Water Mark and Full Table Scans

*Full table scan reads all blocks below the high water mark*



# High Water Mark Impacts

## Should I Worry?

---

- Generally, no
- Natural to have a few empty blocks below the high water mark

## One scenario to monitor

---

- Tables that have a large number of rows inserted and deleted
- Subsequent processing uses many fewer records
- Consider TRUNCATE instead of DELETE

**A data block is “full”**

---

**What does this really mean?**

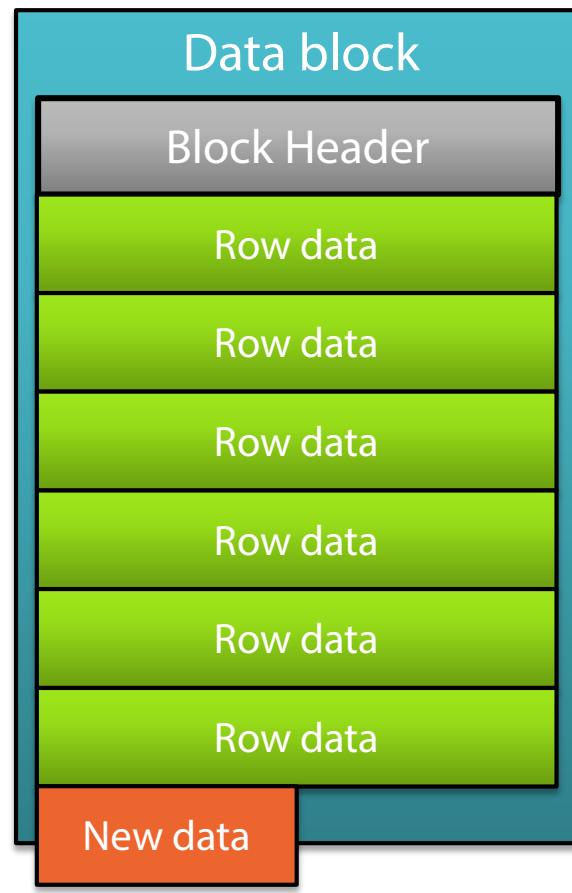
# Storage Implications of Updating Rows

student_id	first_name	last_name	phone	email_address	major_id
208718	Adam	Brady	360-592-7886	AdamBBrady@dayrep.com	3
208723	George	Ward	216-952-0212	GKWard@superrito.com	5
188728	Marcia	Garcia	630-239-4325	MarciaDGarcia@rhyta.com	8
208741	Walter	Stumpf	NULL	WalterCStumpf@einrot.com	12

```
UPDATE students
    SET last_name = 'Anderson'
    WHERE student_id = 188728;
```

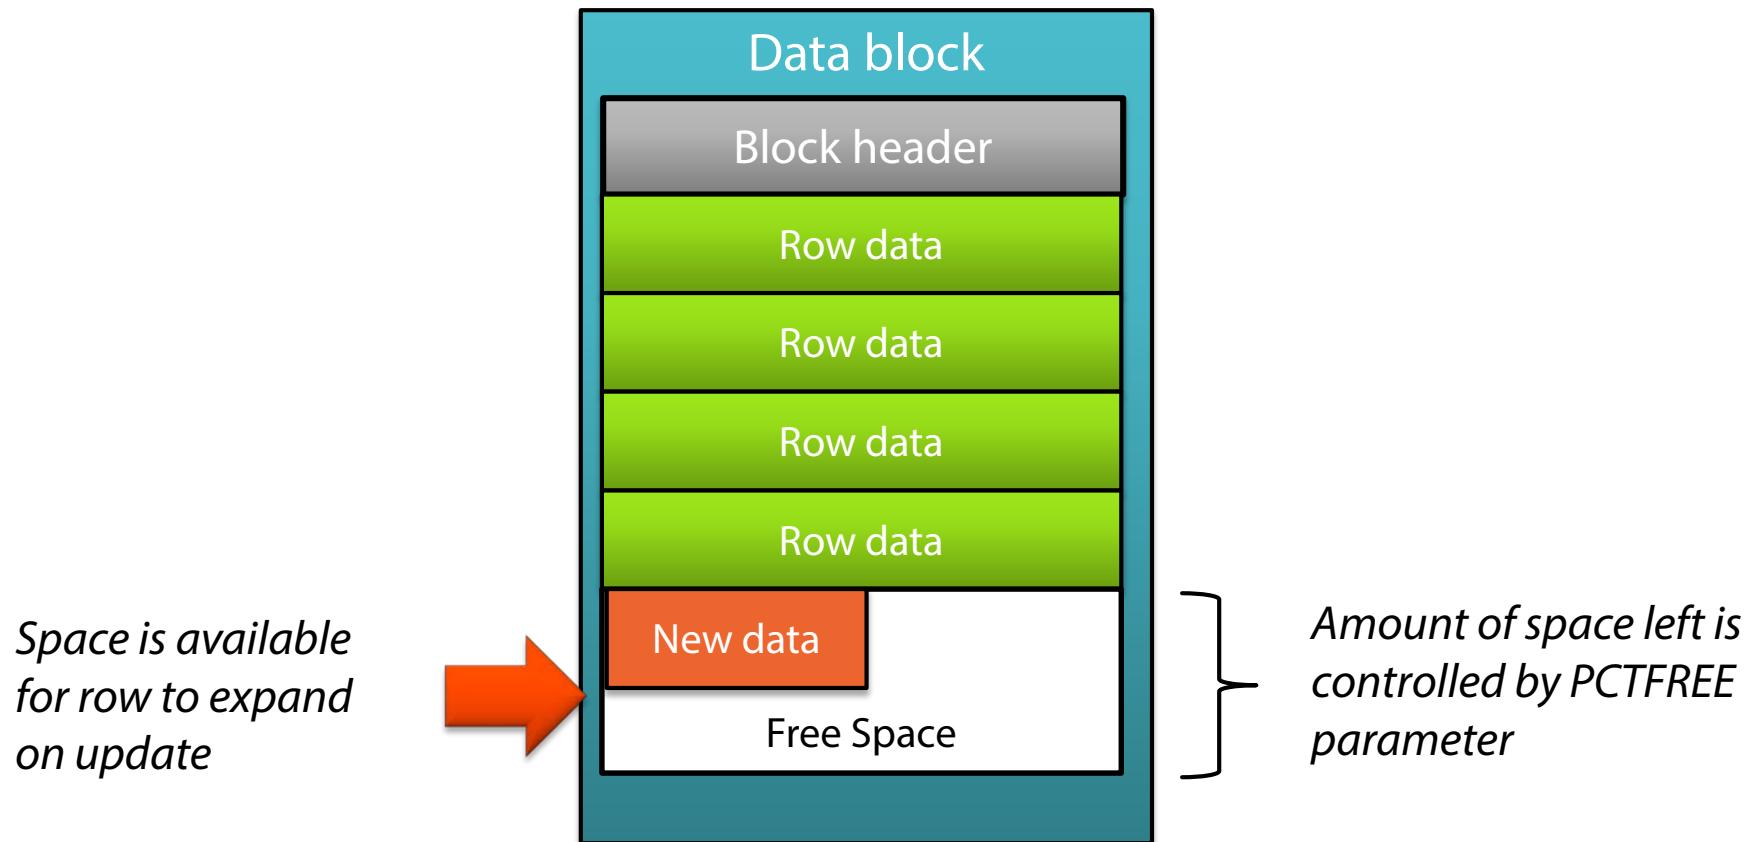
```
UPDATE students
    SET phone = '419-500-5489'
    WHERE student_id = 208741;
```

# PCTFREE Introduction

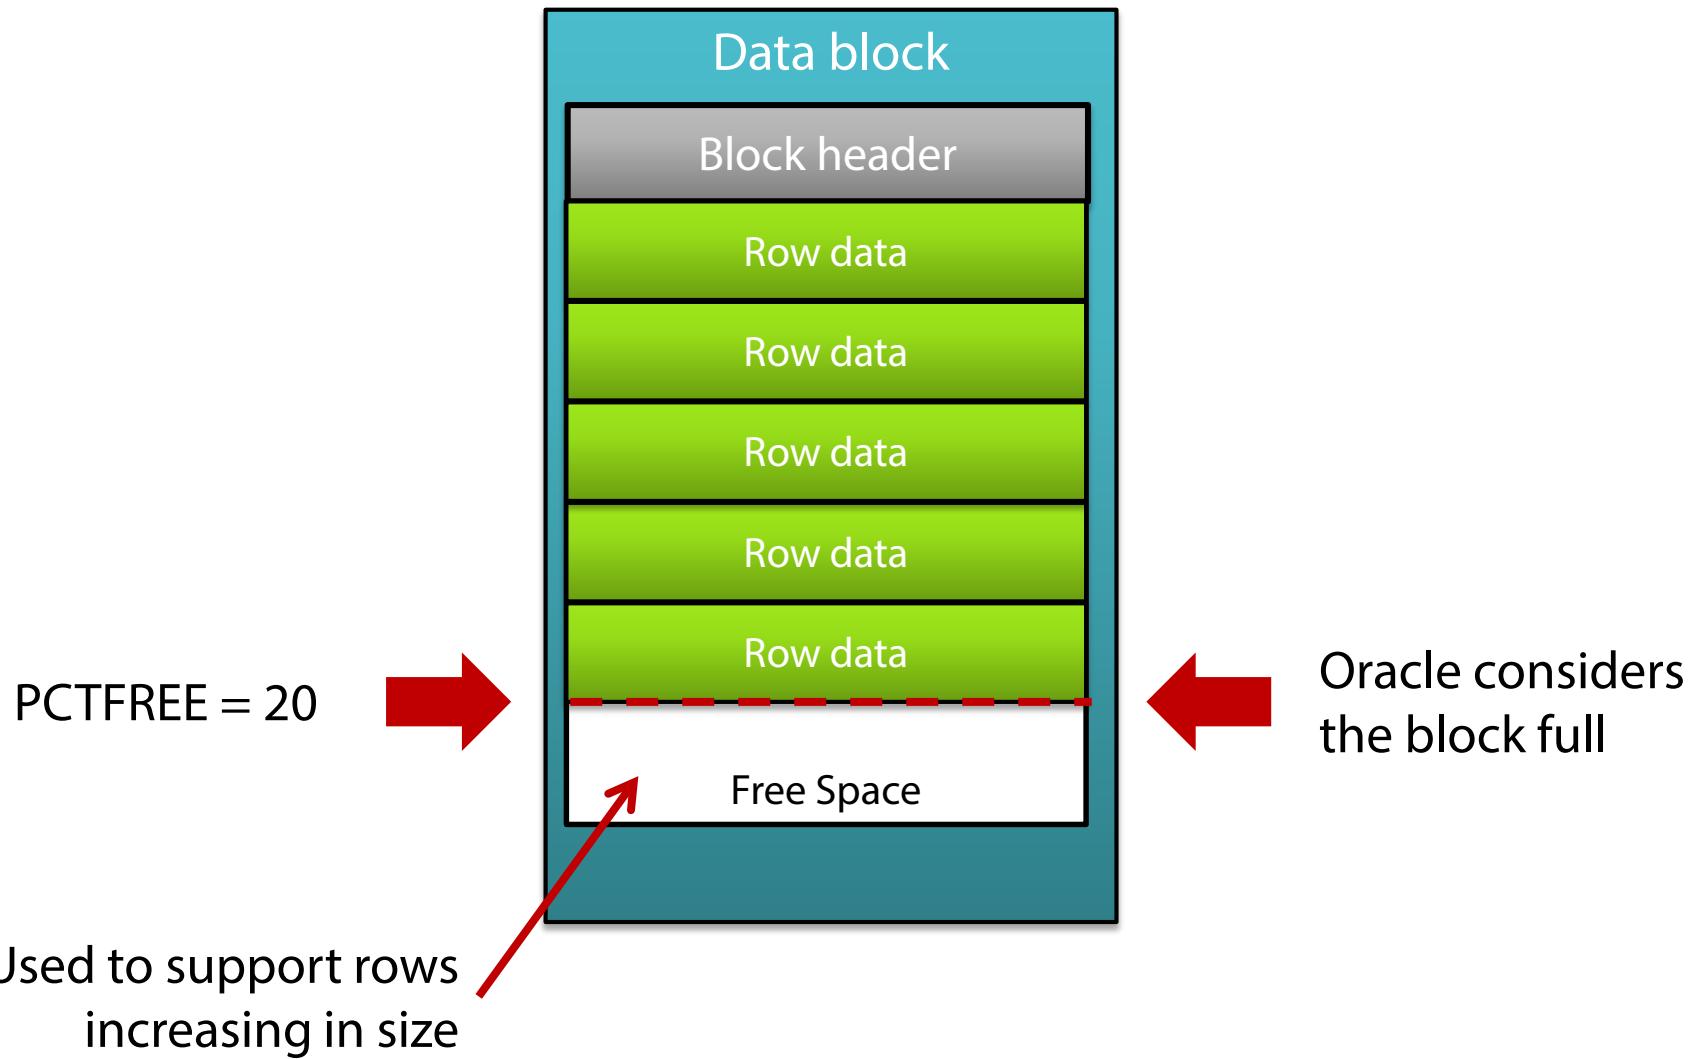


*No space is left for row to expand when the row is updated*

# PCTFREE Introduction



# PCTFREE Example



# Setting PCTFREE

- Default value is 10 percent
- Specify using PCTFREE in storage clause

```
CREATE TABLE students
(
    student_id          NUMBER(10)      NOT NULL,
    first_name          VARCHAR2(40)     NOT NULL,
    middle_name         VARCHAR2(40)     NOT NULL,
    last_name           VARCHAR2(40)     NOT NULL,
    gender              VARCHAR2(1)      NOT NULL,
    street_address      VARCHAR2(50)     NOT NULL,
    city                VARCHAR2(30)     NOT NULL,
    state               VARCHAR2(2)      NOT NULL,
    email               VARCHAR2(80)     NOT NULL,
    ...
)
PCTFREE 20;
```

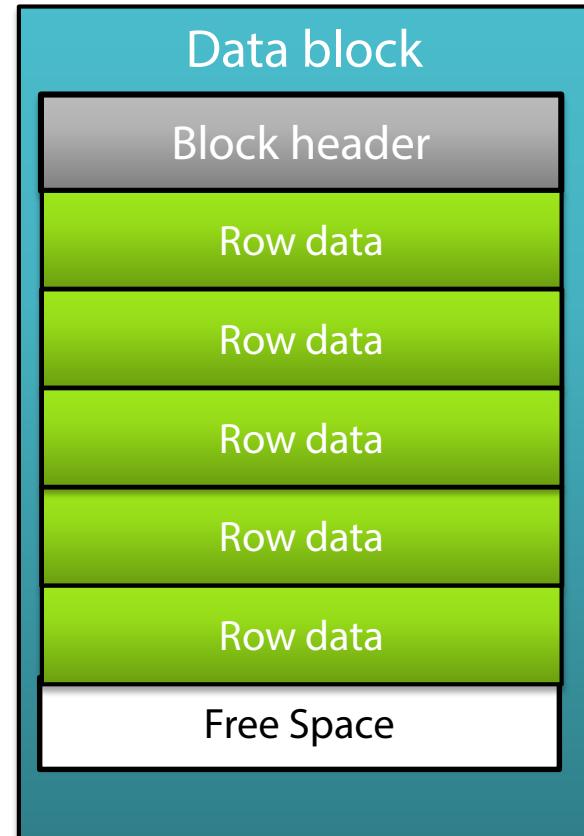
## PCTFREE

- ▶ Applies to both ASSM and MSSM tablespaces
- ▶ Need to understand table data, update patterns to set PCTFREE
- ▶ PCTFREE too low – row migrations will occur on updates
- ▶ PCTFREE too high – space will be wasted in each block

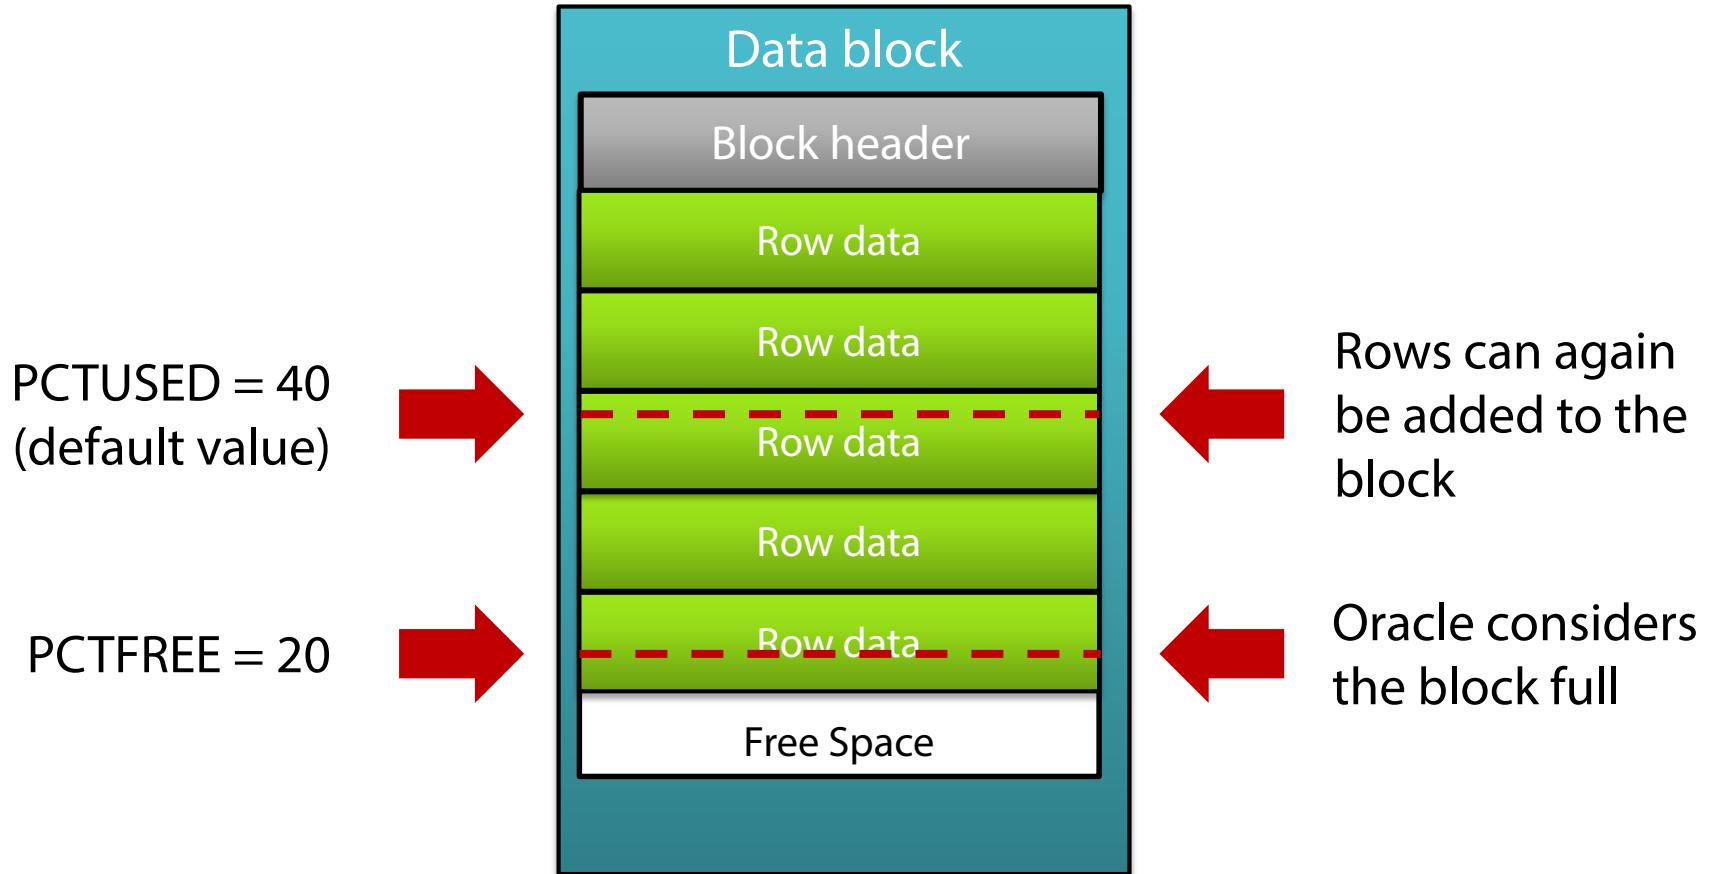
# PCTUSED

Only applies to MSSM tablespaces

Threshold at which Oracle will once again allow inserts into the block



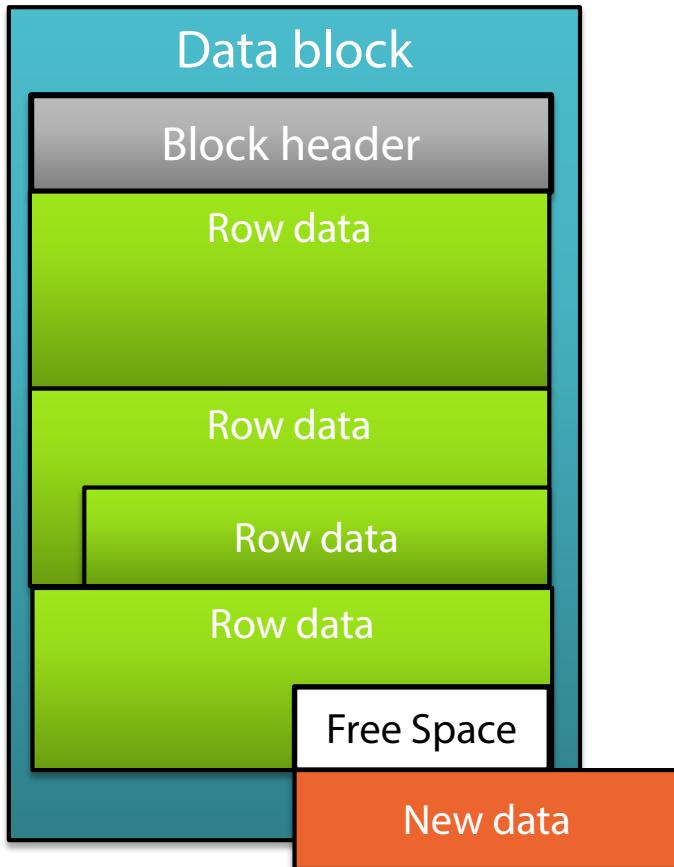
# PCTUSED Example



# Setting PCTUSED

```
CREATE TABLE students
(
    student_id          NUMBER(10)      NOT NULL,
    first_name          VARCHAR2(40)     NOT NULL,
    middle_name         VARCHAR2(40)     NOT NULL,
    last_name           VARCHAR2(40)     NOT NULL,
    gender              VARCHAR2(1)      NOT NULL,
    street_address      VARCHAR2(50)     NOT NULL,
    city                VARCHAR2(30)     NOT NULL,
    state               VARCHAR2(2)      NOT NULL,
    email               VARCHAR2(80)     NOT NULL,
    ...
)
PCTUSED 50;
```

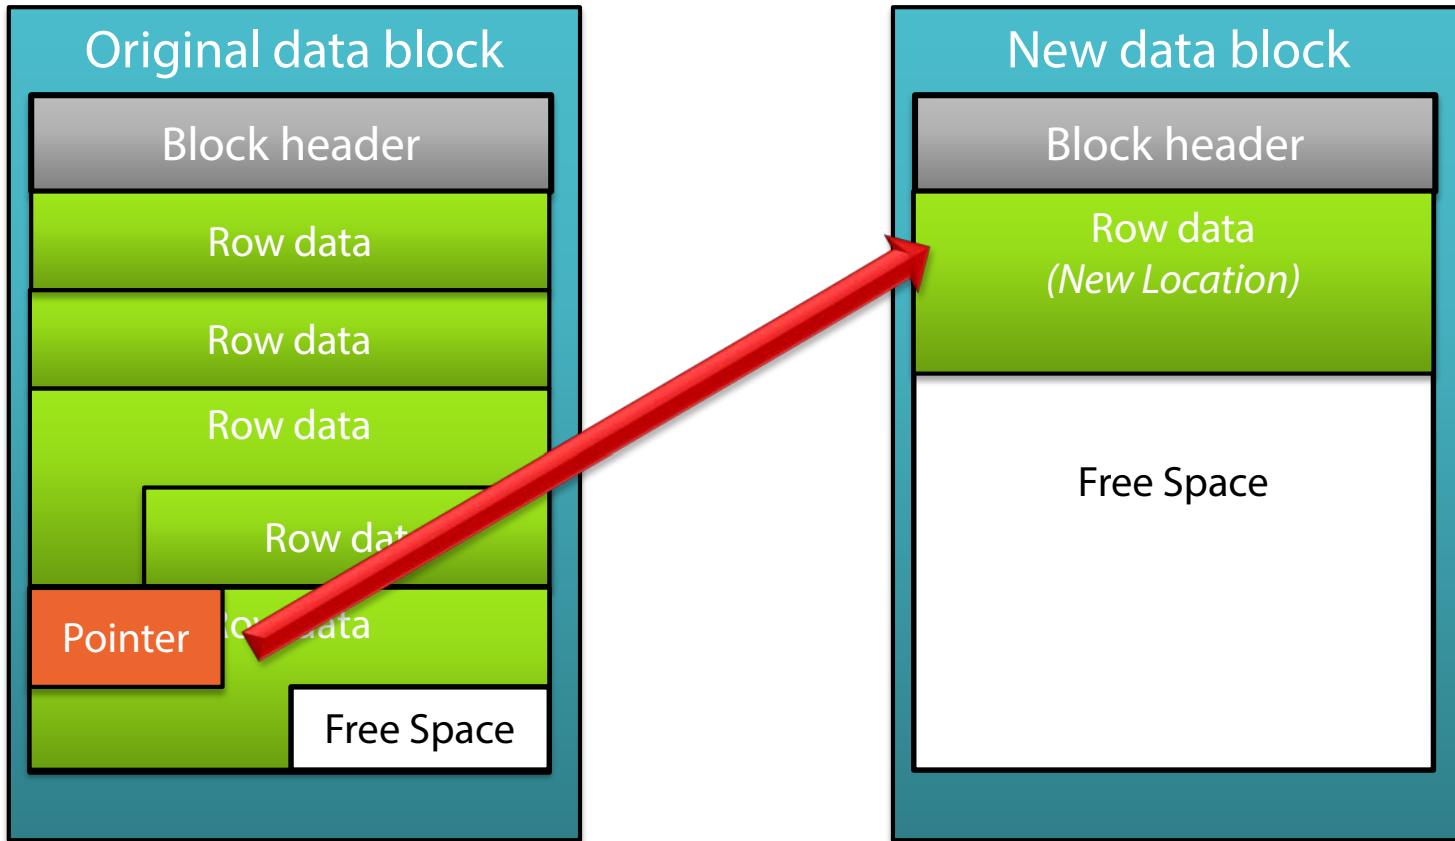
# When Does Row Migration Occur?



Oracle will **migrate** the  
**entire row** to a  
**different** data block

*Updated data will  
not fit in available  
free space of block*

# Row Migration



*Pointer points from  
old location to new  
location*

# **Impacts of Row Migration**

- ▶ Row lookup operations will now require two IO's
- ▶ Full table scans do not incur a penalty
- ▶ Row can be migrated multiple times but Oracle updates the original pointer

# Types of Row Migration

## PCTFREE Too Low

- Insufficient space for rows to expand into
- Many rows will be migrated through normal updates
- Something to be avoided

## Natural Row Migration

- Natural for a small percentage of rows to migrate
- Nothing to worry about

# Detecting Row Migration

**Extra work is coming from row migrations?**

```
SELECT name, value
  FROM v$sysstat
 WHERE name IN ('table fetch continued row',
                 'table fetch by rowid');
```

**Migrated and chained rows per table**

```
SELECT table_name, chain_cnt, num_rows, last_analyzed
  FROM user_tables;
```

# Practical Guidance on PCTFREE

## PCTFREE too low

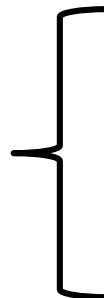
Row migrations will occur

## PCTFREE too high

Waste space with partially filled blocks

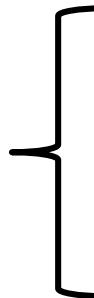
# Using a Low PCTFREE Value

Table Characteristics



- No update activity (inserts only)
- Update will not affect the size of a row

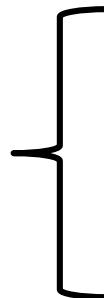
Examples



- Logging tables
- Some data warehouse applications

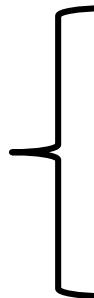
# Using a High PCTFREE Value

Table Characteristics



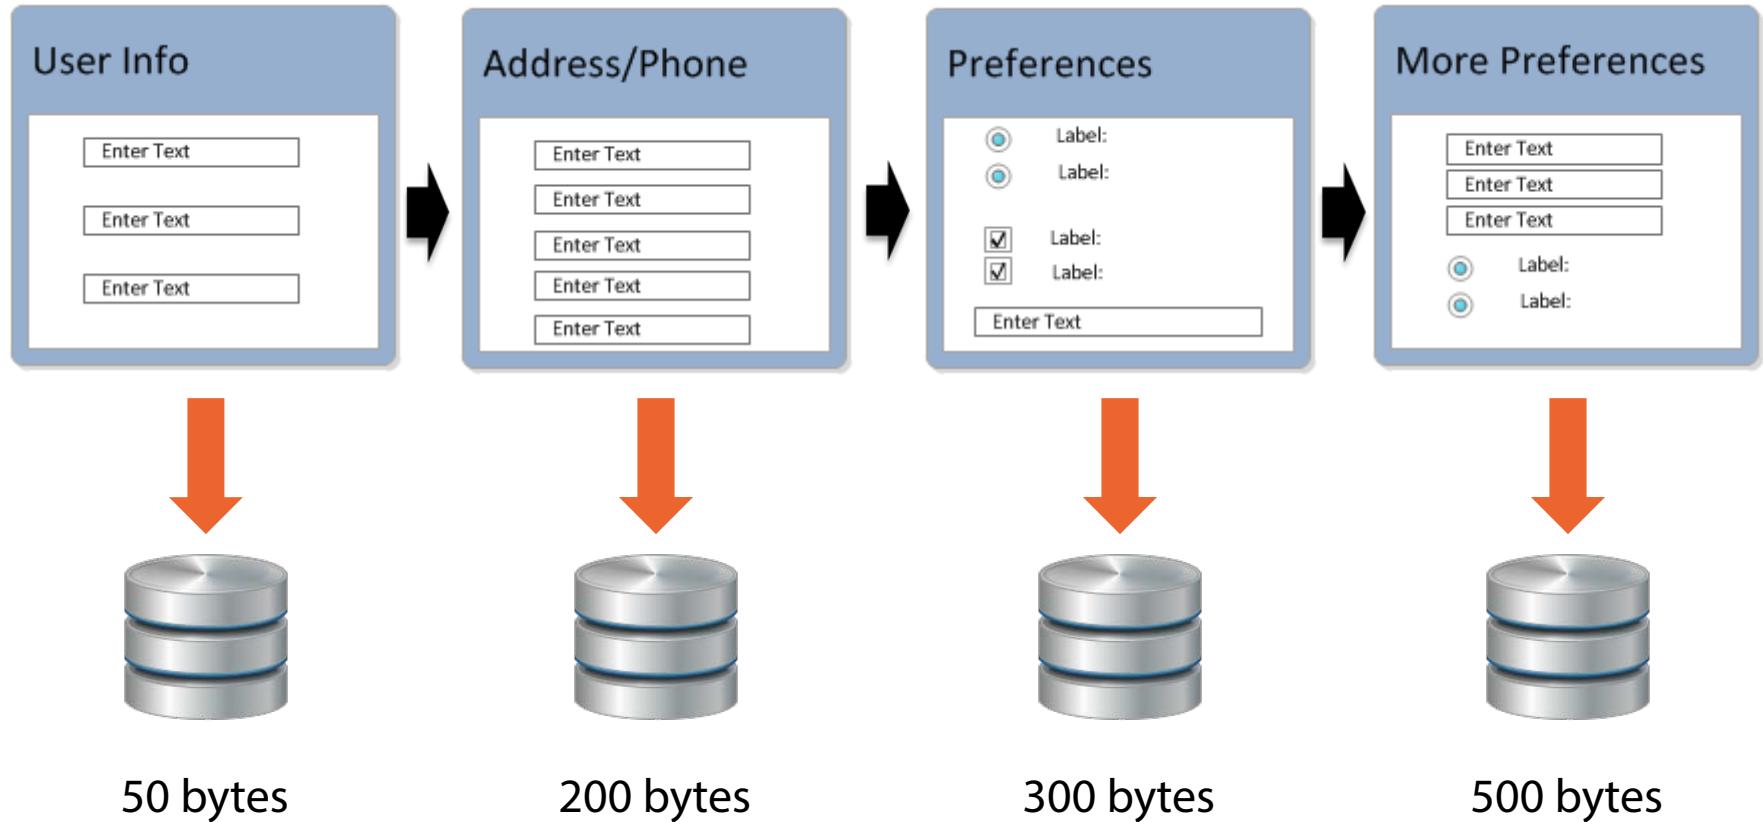
- Data added to rows incrementally
- Updates increase row size

Examples



- Tables seeded with initial value, other data added over time

# A Sample Scenario



# PCTFREE Considerations

Know how your application inserts  
and updates data

Don't be afraid to run a test

# Freelists

## Definition

List of blocks that have **free space** that Oracle can use to **insert new rows** into

## ASSM

Freelists are **automatically** managed  
Major **advantage** of ASSM

## MSSM

**Default** number of freelists is **one**  
Can be a **bottleneck** for concurrent DML

# Setting FREELISTS at Table Creation

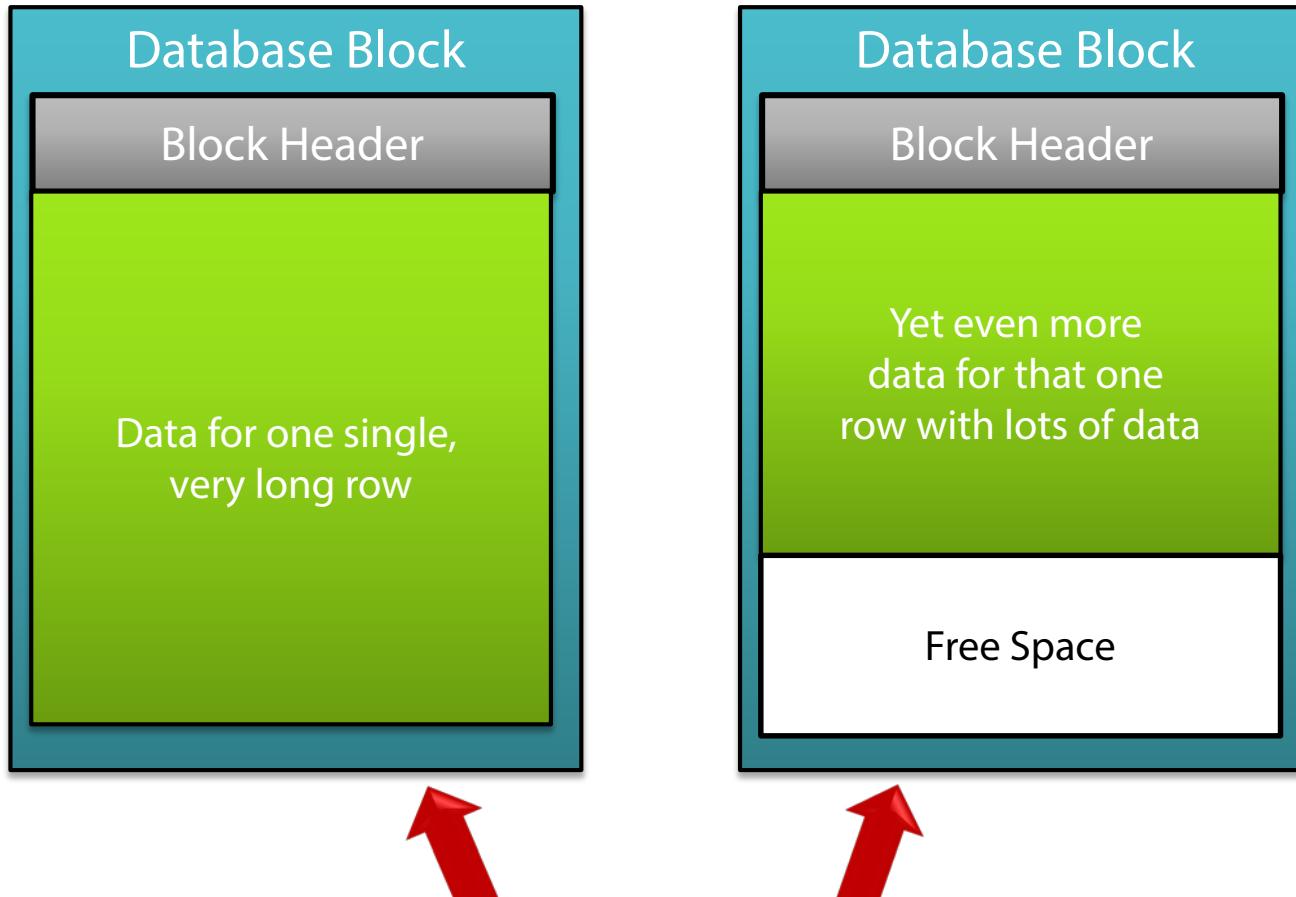
```
CREATE TABLE course_enrollments
(
    course_enrollment_id      NUMBER      NOT NULL,
    course_offering_id        NUMBER(20)   NOT NULL,
    student_id                 NUMBER(10)   NOT NULL,
    grade_code                 VARCHAR2(1)  NULL,
    ...
)
STORAGE (FREELISTS 12);
```

# Changing Freelist on an Existing Table

```
ALTER TABLE course_enrollments  
  STORAGE (FREELISTS 12);
```

- Correct table where freelist was set incorrectly
- Perform a bulk load in parallel

# Row Chaining



*Two IO operations will be required to read this row from disk*

# Table With Row Chaining

## Tablesapce block size

Value between 2 – 32 KB

Assume for example 4 KB

## Sample Table Definition

```
CREATE TABLE products
(
    product_id      NUMBER(8)          NOT NULL,
    sku             VARCHAR2(32)        NOT NULL,
    name            VARCHAR2(128)       NOT NULL,
    description     VARCHAR2(4000)      NOT NULL,
    tech_specs      VARCHAR2(4000)      NOT NULL
);
```

# Detecting Chained Rows

## Tables at risk for row chaining

- Tables with lots of columns (>100 columns)
- Tables with large sized CHAR/VARCHAR2 columns

## Statistics on chained and migrated rows

```
-- Gather statistics on the table
EXEC DBMS_STATS.GATHER_TABLE_STATS(
    ownname => '<<Table Owner Name>>',
    tabname => '<<Table Name>>');

-- Chain count is sum of chained and migrated rows
SELECT table_name, chain_cnt, num_rows, last_analyzed
    FROM user_tables;
```

# Identifying Chained Rows

```
-- Script from $ORACLE_HOME/rdbms/admin/utlchain.sql
CREATE TABLE chained_rows
(
    owner_name          varchar2(128),
    table_name          varchar2(128),
    cluster_name        varchar2(128),
    partition_name      varchar2(128),
    subpartition_name   varchar2(128),
    head_rowid          rowid,
    analyze_timestamp   date
);

-- Analyze the table
ANALYZE TABLE <>table name>> LIST CHAINED ROWS;
```

# Solutions To Row Chaining

## Use a sufficient block size for your table data

---

- ◆ **Estimate** row sizes for each table from **logical data model**
- ◆ Specify the **tablespace block size** based on these estimates
- ◆ **Move** tables with large rows to a **new tablespace**

# Vertically Partition Tables

product_id	sku	name	summary	product_description	technical_specs	product_review
400001	HD-100	256 GB Solid State Hard Drive	....	....	....	....
600483	GC-862	GraphicsStar! Graphics Card	....	....	....	....
7839390	MO-101	27" Flat Screen Monitor – 1080P	.....	....	....	....
8920201	MS-173	Wireless Mouse	.....	....	....	....

# Vertically Partition Tables

product_id	sku	name	summary	product_description
400001	HD-100	256 GB Solid State Hard Drive	....	....
600483	GC-862	GraphicsStar! Graphics Card	....	....
7839390	MO-101	27" Flat Screen Monitor – 1080P	.....	....
8920201	MS-173	Wireless Mouse	.....	....



*Foreign key  
between tables*

product_id	technical_specs	product_review
400001	....	....
600483	....	....
7839390	....	....
8920201	....	....

## Module Summary

Data storage using data blocks

How Oracle accesses data

Tablespaces

PCTFREE

Row Migration

Freelists

Row Chaining

# Oracle Database Concepts Guide

---

<http://bit.ly/Ora12cDatabaseConcepts>

# Managing Tables

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Databases Require Changes Over Time

---

**Adding columns**

**Changing data types**

**Renaming objects**

# Module Contents

Oracle  
Data  
Dictionary

Database  
Statistics

Managing  
Tables

Managing  
Columns

# Oracle Data Dictionary

Series of **views** containing metadata about all **database objects** in Oracle

# Data Dictionary Course Content Focus

## Graphical presentation varies by tool

Become familiar with feature set of tool you use

## Be aware of underlying views

Use for writing custom queries, customizing output format

# Three Sets of Views

**dba\_\***

Information about all objects in the database

**user\_\***

Information about objects owned by the current user

**all\_\***

Information about all objects the current user has permission to

# USER\_OBJECTS View

## *Definition*

Contains	An entry for every database object (table, view, index, stored procedure, etc...)
Oracle Documentation	<a href="http://bit.ly/Ora12cAllObjects">http://bit.ly/Ora12cAllObjects</a>

## *Key columns*

object_name	Name of the table, view, index, etc...
object_type	Indicated if this object is a table, view, etc...
created	Timestamp of when this object was created
last_ddl_time	Timestamp of when the definition of this object was last changed
status	Indicated if the object is valid or invalid

# **USER\_TABLES View**

## *Definition*

Contains	One row for each table
Oracle Documentation	<a href="http://bit.ly/Ora12cAllTables">http://bit.ly/Ora12cAllTables</a>

## *Key columns*

table_name	Name of the table
num_rows	Number of rows in the table
blocks	Number of blocks that have been used in the table
avg_row_len	Average length in bytes of a row in the table
last_analyzed	Date that stats were most recently gathered against the table

# **USER\_TAB\_COLS View**

## *Definition*

Contains	An entry for every column in a table or a view
Oracle Documentation	<a href="http://bit.ly/Ora12cAllTabCols">http://bit.ly/Ora12cAllTabCols</a>

## *Key columns*

table_name	Name of the table or view the column belongs to
column_name	Name of the column
data_type	Data type of this column
num_distinct	Number of distinct values in this column
num_nulls	Number of rows in this column with a null value

# **USER\_CONSTRAINTS View**

## *Definition*

Contains	An entry for each constraint (primary key, foreign key, check constraint)
Oracle Documentation	<a href="http://bit.ly/Ora12cAllConstraints">http://bit.ly/Ora12cAllConstraints</a>

## *Key columns*

table_name	Name of the table the constraint is associated with
constraint_name	Name of the constraint
constraint_type	Code that indicates the type of the constraint
r_constraint_name	For foreign keys, the name of the unique constraint in the parent table
deferrable	Indicates if the constraint can be deferred

# **USER\_INDEXES View**

## *Definition*

Contains	An entry for each index
Oracle Documentation	<a href="http://bit.ly/Ora12cAllindexes">http://bit.ly/Ora12cAllindexes</a>

## *Key columns*

table_name	Name of the table the index is on
index_name	Name of index
uniqueness	Indicated if this is a unique index
distinct_keys	Number of distinct keys in the index
last_analyzed	Date that stats were most recently gathered against the index

# **USER\_IND\_COLUMNS View**

## *Definition*

Contains	An entry for each column in an index
Oracle Documentation	<a href="http://bit.ly/Ora12cAllIndColumns">http://bit.ly/Ora12cAllIndColumns</a>

## *Key columns*

table_name	Name of the table the index is on
index_name	Name of index
column_name	Name of the column
column_position	Position of the column in the index

A massive iceberg, mostly white with some blue at the base, sits in a dark blue ocean. The sky above is a clear, pale blue with a few wispy clouds.

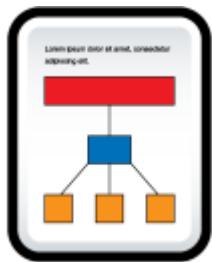
# Data Dictionary Summary

Invest some time in learning what is available

Most commonly used views

- <http://bit.ly/Ora12cDataDictionaryViews>

# Life of a SQL Statement



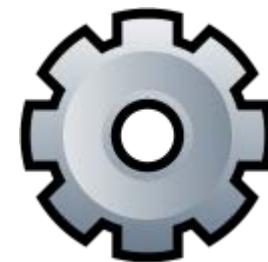
## Parse Phase

Check SQL Syntax  
Check Object Permissions  
Query Rewrite Process



## Query Optimization

Evaluate Statistics  
Create Execution Plan



## Execution Phase

Read Blocks  
Filter Rows  
Sort Data

# How Are Database Statistics Used?

By the Oracle Optimizer to determine the most efficient plan for your SQL statement

To estimate how much processing is required for each option

# How Are Database Statistics Gathered?

## Automatically

- Scheduled by Oracle
- Oracle detects objects with significant changes

## Manually

- Executed by user
- Used after large insert/delete operation
- Useful in test environments

# Gathering Statistics Manually

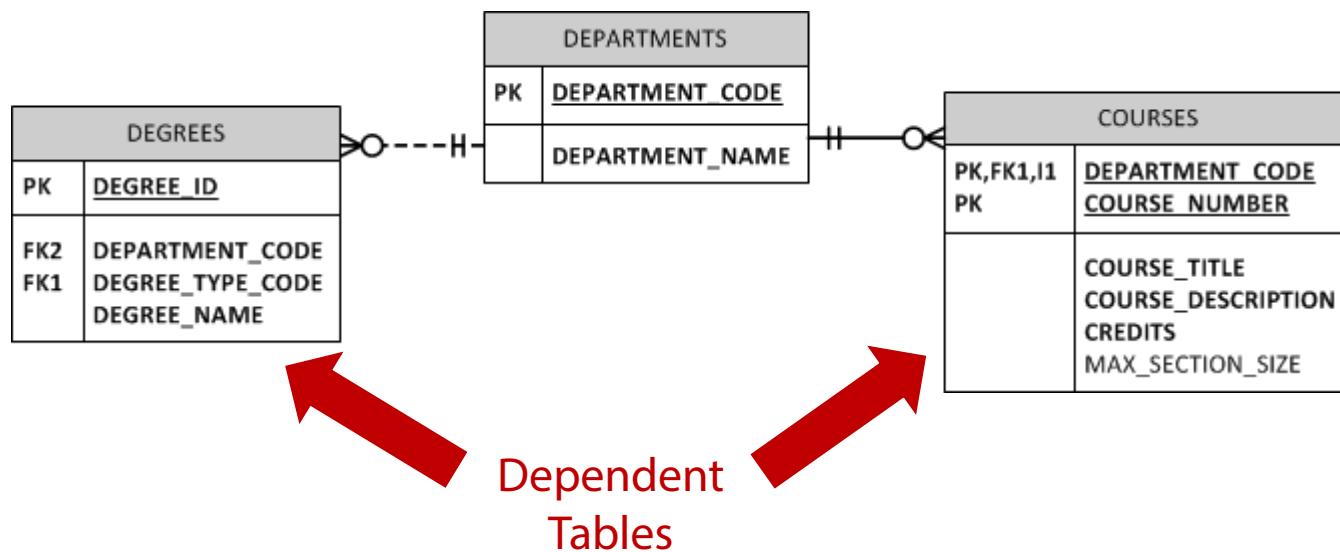
```
-- Update for one table and all indexes on table
EXEC DBMS_STATS.GATHER_TABLE_STATS(
    ownname => '<<Table Owner Name>>',
    tabname => '<<Table Name>>',
    estimate_percent => <<percent>>);

-- Update for all objects in a schema
EXEC DBMS_STATS. GATHER_SCHEMA_STATS(
    ownname => '<<Schema Owner Name>>',
    estimate_percent => <<percent>> );
```

<http://bit.ly/Orac12cDbmsStats>

# Dropping a Table

```
-- Basic DROP TABLE syntax  
DROP TABLE departments;
```



# Foreign Keys and Dropping a Table

```
-- Drop dependent tables first
DROP TABLE degrees;
DROP TABLE courses;

-- Now drop the table of interest
DROP TABLE departments;
```

- **Problematic if the dependent tables contain more dependencies**
- **We don't want to drop a significant portion of the database**

# Drop Table With Cascade Constraints

```
-- Drop table and foreign keys that reference this table  
DROP TABLE departments CASCADE CONSTRAINTS;
```

- **Dependent tables stay intact**
  - Only their foreign keys are dropped
- **Possible to recreate the table, re-add foreign keys**

# Renaming a Table

```
-- Syntax to rename an object in Oracle  
-- RENAME old_name TO new_name;  
  
RENAME departments TO depts;
```

- **Oracle will automatically transfer all constraints and indexes**
- **Views, stored procedures, functions, synonyms will become invalid**
  - Find invalid objects using all\_objects view
  - You need to manually fix these

# Renaming a Table

```
-- Syntax to rename an object in Oracle  
-- RENAME old_name TO new_name;  
  
RENAME departments TO depts;
```

- **Oracle will automatically transfer all constraints and indexes**
- **Views, stored procedures, functions, synonyms will become invalid**
  - Find invalid objects using all\_objects view
  - You need to manually fix these

# Making a Table Read Only

## Place a table in read only mode

```
ALTER TABLE table_name READ ONLY;
```

- No DML allowed (INSERT, UPDATE, DELETE)
- TRUNCATE not allowed
- Cannot add, modify, remove columns

## Return a table to read write mode

```
ALTER TABLE table_name READ WRITE;
```

## Adding Columns to Tables

---

**Frequent** need over the lifetime of a database

Required capability to respond to **changing business needs**

# Adding Columns to a Table

```
ALTER TABLE students ADD
(
    immunization_form_received    VARCHAR2(1) DEFAULT 'N' NOT NULL,
    immunization_form_date        DATE          NULL
);
```

# Adding Columns With a Default Value

```
ALTER TABLE students ADD
(
    immunization_form_received    VARCHAR2(1) DEFAULT 'N' NOT NULL,
    immunization_form_date        DATE         NULL
);
```

- Oracle substitutes the default value at query time for existing rows
- Existing rows do not have to be updated

# Adding Columns – Breaking Insert Statements

Insert statement without column names

```
INSERT INTO courses  
VALUES ('CS', '401', 'Operating Systems', 4.0);
```

Solution: add the column as invisible

```
ALTER TABLE courses ADD  
(  
    textbook_name    VARCHAR2(50)    INVISIBLE NULL  
);
```

# **Drop a Column From a Table**

Logical Delete

Physical Delete

# **Physical Delete of Column**

```
ALTER TABLE table_name  
DROP COLUMN (column_name1, column_name2);
```

# Physical Delete of Column



Views, procedures will be invalidated

Your responsibility to fix these



Oracle will update each block in table

Can be a performance impact on a large table

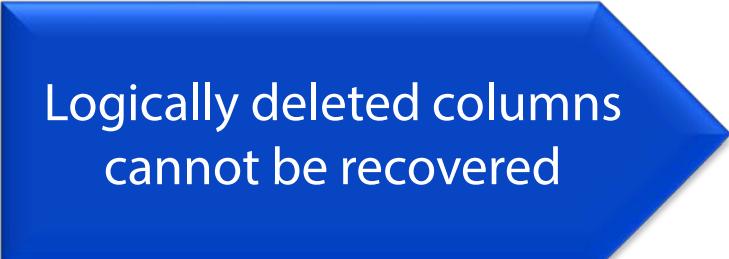
# Logical Delete of Column

```
ALTER TABLE table_name  
    SET UNUSED (column_name1, column_name2);
```

*At some later time:*

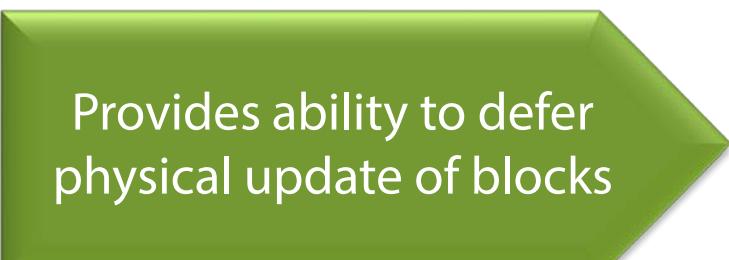
```
ALTER TABLE table_name  
    DROP UNUSED COLUMNS;
```

# **Logical Delete of Column**



Logically deleted columns  
cannot be recovered

Regard this as a one way trip



Provides ability to defer  
physical update of blocks

Can occur during off hours or a  
maintenance window

# Dropping Columns and Key Constraints

- Attempt to drop column in a primary or foreign key will result in an error
- To drop the column and the primary/foreign key use

```
ALTER TABLE table_name  
    SET UNUSED (column_name1, column_name2)  
    CASCADE CONSTRAINTS;
```

- Is dropping a column in a key constraint what you really want?

## Changing a Column Data Type

---

**Possible** to do, but there are **limitations**

# Changing a Column Data Type

Operation	Examples	Allowed?
<b>Increasing size of data type</b>	VARCHAR2(30) → VARCHAR2(50) NUMBER(8,4) → NUMBER(12,6)	Allowed
<b>Decreasing size of VARCHAR</b>	VARCHAR2(50) → VARCHAR2(25)	Allowed if no data truncated
<b>Decrease precision, scale of NUMBER</b>	NUMBER(8,4) → NUMBER(12,6)	Allowed if column empty
<b>Change to different data type</b>	NUMBER(8,4) → VARCHAR2(10) VARCHAR2(10) → DATE	Allowed if column empty

# Change Column Data Type Syntax

```
ALTER TABLE table_name MODIFY
(
    column_name      new_data_type
);
```

# Module Summary

## Oracle Data Dictionary

What is the data dictionary  
Useful views

## Database Statistics

What are database statistics  
How are they used  
How are stats gathered

## Managing Tables

Dropping tables  
Renaming tables

## Managing Columns

Adding columns  
Renaming columns  
Changing column data types

# Specialized Table Types

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Module Outline

Temporary Tables

Hold temporary results during a transaction or session

External Tables

Easily import data from flat files

## Temporary Tables Overview

---

- ▶ **Defined statically**
- ▶ **Table definition available to all sessions**
- ▶ **Data is private to a session**
- ▶ **Act as a private workspace**

# Basic Create Table Syntax

```
CREATE GLOBAL TEMPORARY TABLE temporary_courses
(
    department_code      VARCHAR2(2)      NOT NULL,
    course_number        NUMBER(3,0)       NOT NULL,
    course_title         VARCHAR2(64)      NOT NULL,
    course_description   VARCHAR2(512)     NOT NULL,
    credits              NUMBER(3,1)       NOT NULL,
    CONSTRAINT pk_courses PRIMARY KEY
        (department_code, course_number),
)
ON COMMIT DELETE ROWS
-- ON COMMIT PRESERVE ROWS;
```

# Preserve Rows Options

## **ON COMMIT DELETE ROWS**

---

Rows automatically deleted when a transaction commits  
Default if no option is specified

## **ON COMMIT PRESERVE ROWS**

---

Rows preserved when a transaction commits  
Rows will be deleted when session ends

# Additional Rules

**Temporary tables cannot participate in foreign key relationships**

**Indexes can be created on temporary tables**

**Temporary tables do not have any statistics by default**

**Database stats are per session starting in Oracle 12c**

- Myles, there will be a demo in here, so it isn't like we are going between two slides in the same format
- This slide obviously won't be there either (unless you *really* think I should leave it in.)

# External Tables

**Common need: importing a flat file**

**Allow flat file data to be processed by SQL**

**Data can be filtered, joined or inserted into other tables**

# External Table Properties

## Import only

Only used for **importing** data, not exporting data

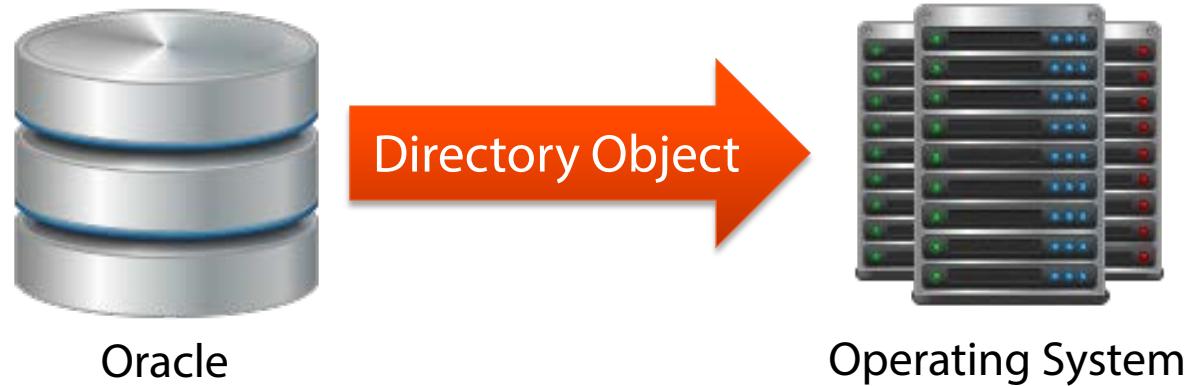
## File formats

Import both **fixed width** and **delimited** files

## File location

Files must be **located** on the **Oracle server**

# Oracle and Directory Objects



- **Directory objects tell Oracle where to find data on the server**
- **Will need system administrator and DBA permissions to set up**

# Steps to Create a Data Directory

Create directory in operating system



Assign operating system permissions



Create directory object in Oracle



Give Oracle users read access to the directory object

# Oracle Directory Objects

```
-- Create the directory object  
CREATE DIRECTORY data_import AS  
AS 'C:\data_import\';  
  
-- Grant permissions so a user can use  
GRANT READ ON DIRECTORY data_import  
TO student;
```

## **Create one directory per schema/application**

---

-  **Protect sensitive data**
-  **Clean separation between applications**

# External Table Resources

## Oracle resources

---

External Tables Concepts - <http://bit.ly/Ora12cExternalTables>

Oracle Loader Access Driver - <http://bit.ly/OracleLoaderAccessDriver>

## Sample external table definitions

---

On my blog - <http://bit.ly/SampleExternalTables>

## When you have errors

---

Don't forget about the log and discard files

# **Specialized Tables Summary**



Temporary  
Tables



External Tables

# Indexes

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

## **Tables**

---

Store data for your application

## **Indexes**

---

Enable rapid searching of your data

# Module Outline

What is an index?

Creating standard  
(B-tree) indexes

Unique indexes

Function based  
indexes

Bitmap indexes

Indexing tips

# Scanning a Table for Matching Rows

ID	LAST_NAME	FIRST_NAME
213458	Marshall	George
134876	Smith	Ann
142782	Lewis	Gregory
152212	Baker	Charles
108759	Schultz	Sara
254785	Nguyen	Tom
103897	Rodriguez	Carlos
204872	Anderson	Adam
198344	Lin	Jason
175485	Chao	Kiho
154852	Wilson	Jeff

## For small tables

---

- Few rows to search, so data found quickly

## For large tables

---

- Takes a long time – lots of rows to read and process
- Computationally intensive to read and process every row

An index is a **separate data structure** designed to  
**speed up** the retrieval of rows from a table

The **most common** type of index in Oracle is based is a  
**B-tree (balanced tree) index**

# Index Key

## Definition

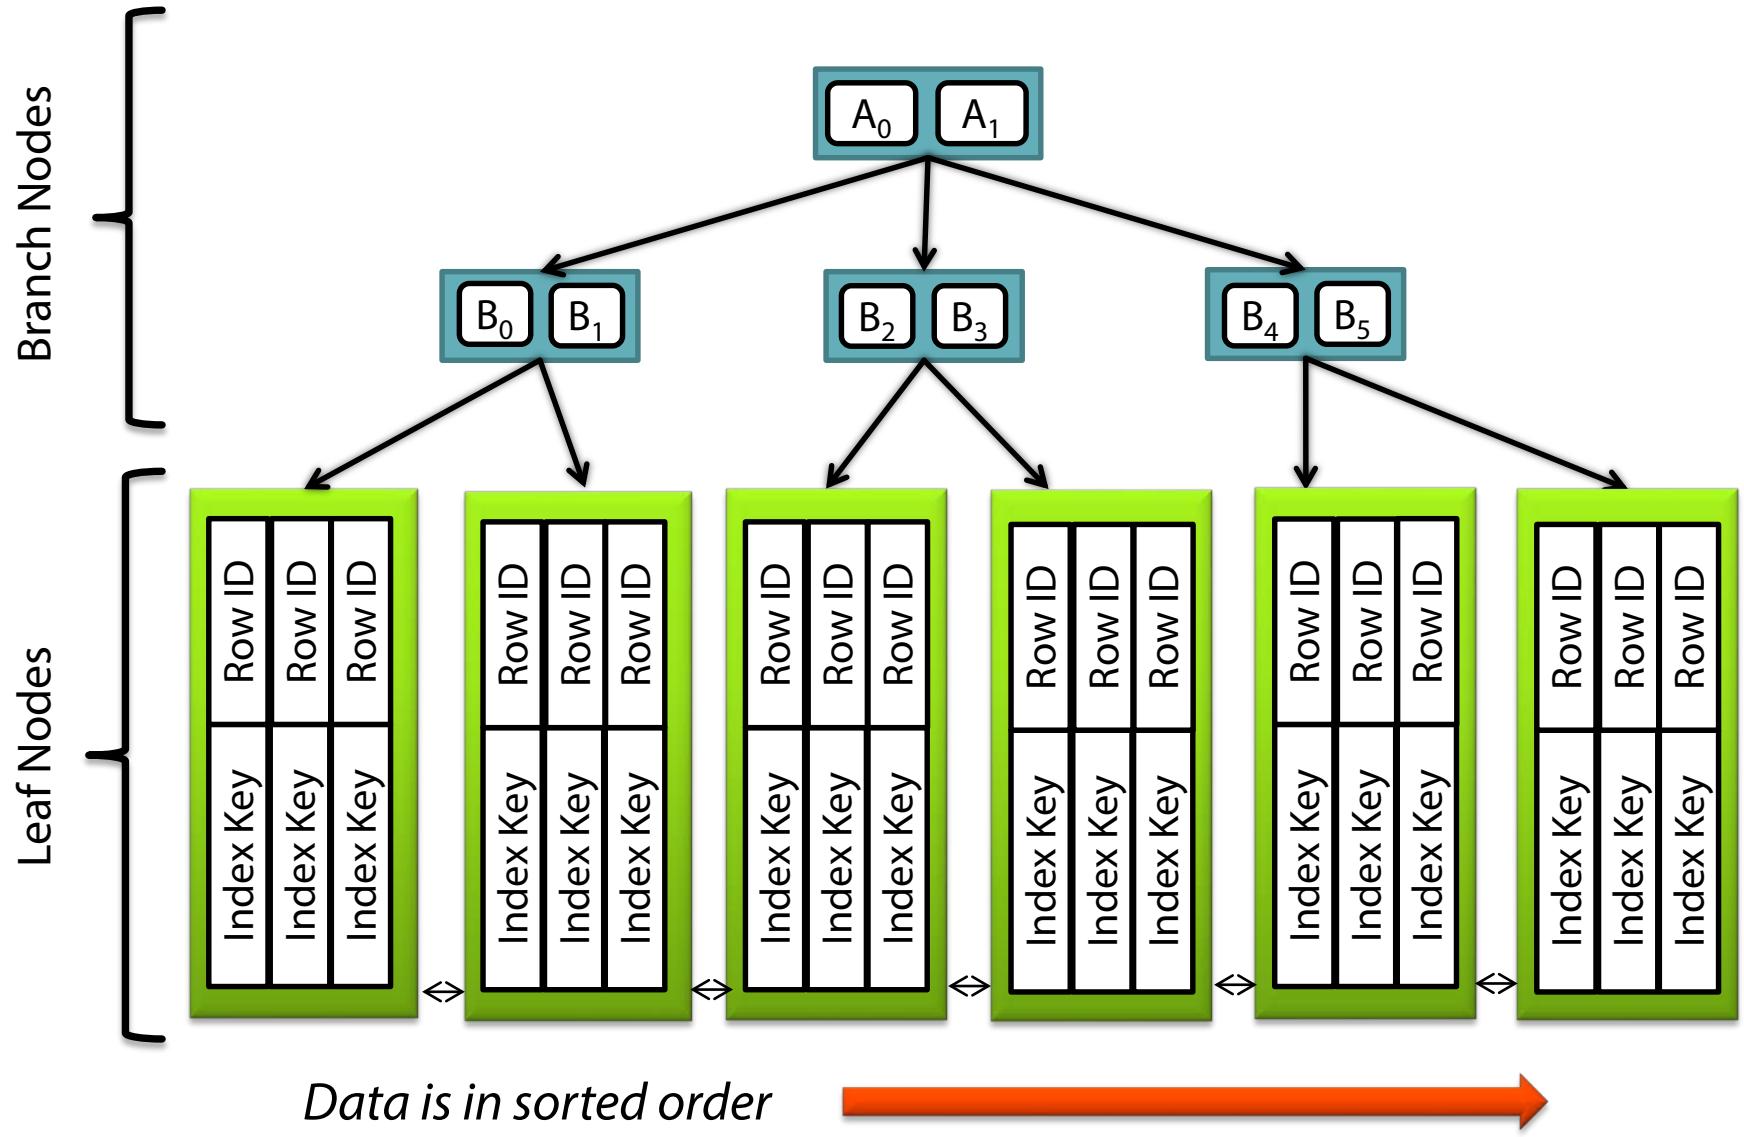
A subset of columns used to organize and search the index

## Phone book example

Last name and first name would be the index key

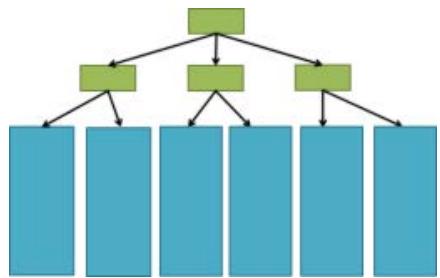
Francis & Diane	N7971 Pigeon Rd Sherwood	920 989-3175	727-4403
Lee	2624 N Owaisa St Aptn	739-5098	Louis
PLANK	Dave & Carol	245 E Bell Nenh	725-2383
David	8928 Pine La Fremont	920 446-3940	PLOETZ
David	3119 Selma Ct Aptn	882-9223	Ge
David & Kathleen	3533 N Winterset Dr Aptn	997-9810	Keith 26
G Neenah		729-0009	Pahl 303
Gene F	7791 Haase Rd Larsen	920 836-3217	Scott W
J	2102 Gateway Pl Nenh	720-9697	PLOG
Joe & Cheree	Appleton	734-6713	Chr
Joseph	Appleton	993-1886	PLONSKE
Mark A	2975 W Lawrence Aptn	997-0349	PLONT M
Randy & Connie		7783 Boom Bay Heights Rd Larsen	PLOOR A
Richard	624 Milw Mnsh	920 836-2920	Edmu
Robt & Phyllis	408 Hawthorne Nenh	720-0043	Jame
Tom	2840 W 1st Ave Aptn	729-1697	PLOTE B
Wm J Jr	7 Reid Ct Aptn	738-6710	Brice
PLANKEY	Kall	739-2947	PLOTZ
PLANNER	Bill & Bonne	23 N Greves Ct Grnd Cht	Chris
Dorothy	1029 W Windtree Dr Aptn	731-2674	Mari
PLANTE	Andrew	833 Congress St Nenh	Mav
Dennis & Catherine	2303 Marathon Ave Nenh	687-0293	PLOUC
J	600 E Wilson Av Aptn	558-4567	PLOU
Michael	119 Olde School Rd Nenh	725-9194	604
PLANTIKO	Kurt & Melissa	739-2811	Ma
PLANTIKOW	James E	1901 E Ashbury Dr Aptn	PLUC
Jerome	410 W Henry Kauk	993-8357	PLUE
PLANTZ	Earl John	213 S Lincoln Kmbr	PLUI
PLASKI	Thomas W & Tammy B	766-3520	WS
W	8716 Spring Rd Hrtvnl	788-3806	PLU
PLASS	Donald	729-6398	D
Earl	415 E Second Kmbr	734-4279	D
R	322 E Third Kmbr	779-0359	PL
R	1500 N Appleton Aptn	687-1117	D
Tim	129 Mary Cmb Lks	788-5241	PL
Troy	W 3569 Mullen Rd Symr	788-3920	P
PLASTER	Cherie	788-0618	P
PLATE	Allan	731-7433	P
Anita	49 1st St Hilbert	731-0046	P
Arlyn	13 A S 6th Hilbert	920 853-3193	P
David & Charity	3675 N Terri La Aptn	920 853-3390	P
Ervin	445 Ridgeway Dr Brillion	920 853-3300	P
Kenneth	32 W Main Hilbert	920 853-3768	F
Kevin	N7240 Irish Rd	920 853-3863	F
Kevin & Laurie	42 Creek St Hilbert	920 853-3330	F
W2920	Faro Springs Rd New Holstein	920 853-3768	F
Kim & Katie	N-4392 Weeks Rd Chilton	920 853-3972	F
		920 853-3104	
		920 849-8373	

# Standard (B-Tree) Index



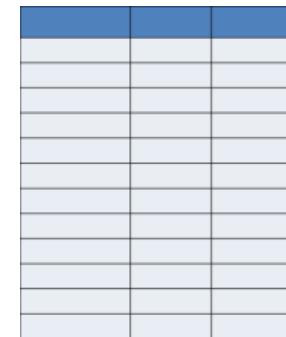
# Using an Index in a SQL Statement

**Index**

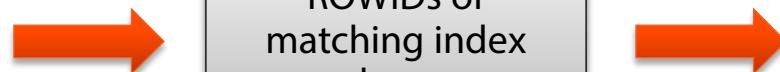


**Index  
lookup  
operation**

**Table**

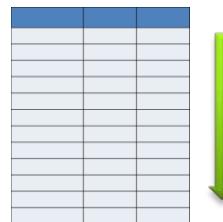


**Table access  
by ROWID**

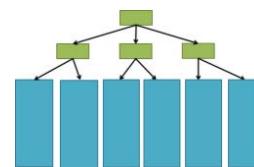


# Searching an Index vs. a Full Table Scan

Table Scan



Index (B-Tree)



Search efficiency

$O(n)$

$O(\log n)$

Comparisons for searching 1 million items

1,000,000

19

## Why Index?

Necessary on larger tables to get **good performance**

SQL statements will **complete faster**

Statements will use **fewer resources** on the Oracle server

# Defining an Index

```
CREATE INDEX <<index name>>
    ON <<table name>>
        (<<column 1 name>>, <<column 2 name>> ...)
TABLESPACE <<index tablespace name>>;
```

# Defining an Index

```
CREATE INDEX <<index name>>
  ON <<table name>>
    (<<column 1 name>>, <<column 2 name>> ...)
  ONLINE
  TABLESPACE <<index tablespace name>>;
```

- ◆ Only available in Oracle Enterprise Edition
- ◆ Oracle will not lock the table for DML during index creation

# Unique Index Introduction

## Non-unique indexes

Multiple rows can have the **same value** for the columns in the **index key**

Example: first and last name

## Unique indexes

Every row must have a **unique value** for the columns in the **index key**

Example: email address

# Defining a Unique Index

```
CREATE UNIQUE INDEX <<index name>>
    ON <<table name>>
        (<<column 1 name>>, <<column 2 name>> ...)
    [ONLINE]
    [TABLESPACE <<index tablespace name>>];
```

# Function Based Index Introduction

## Standard index

Index is created over the **actual column value**

## Function based index

Index is created over a **derived value** returned by a function

# Classic Example: Case Insensitive Search

```
SELECT *
  FROM applications
 WHERE last_name = 'Thompson'
   AND first_name = 'Eric';
```

**What about using the UPPER() function to remove case sensitivity?**

```
SELECT *
  FROM applications
 WHERE UPPER(last_name) = UPPER('Thompson')
   AND UPPER(first_name) = UPPER('Eric');
```

***PROBLEM! Oracle will no longer use an index on last\_name, first\_name columns!***

# Creating a Function Based Index

```
CREATE INDEX fx_applications_fullname  
ON (applications)  
( UPPER(last_name), UPPER(first_name) );
```

**SQL statement with UPPER() function can use this index**

```
SELECT *  
FROM applications  
WHERE UPPER(last_name) = UPPER('Thompson')  
AND UPPER(first_name) = UPPER('Eric');
```

# Function Based Index Rules

- ◆ Can use built in Oracle or user defined functions
- ◆ Functions must be deterministic
- ◆ Indexes can mix normal column values and derived values from functions

# Bitmap Index Introduction

Intended to solve a very **different problem** than a B-tree index

Targeted at columns with **low number** of **distinct values**

Only suitable for **data warehouse** or **read only** applications

# Columns With Few Distinct Values

*Table data*

Customer ID	Gender	Married	Age Range
1001	F	N	18-34
1002	F	Y	35-49
1003	M	Y	35-49
1004	F	Y	35-49
1005	F	Y	18-34
1006	M	Y	62+
1007	F	N	50-61
1008	M	Y	35-49
1009	F	Y	35-49
1010	M	N	35-49
1011	M	Y	35-49
1012	F	N	18-34



**Few distinct values in each column**

# Bitmap Indexes – Simple Example

Table Data

Customer ID	Gender
1001	F
1002	F
1003	M
1004	F
1005	F
1006	M
1007	F
1008	M
1009	F
1010	M
1011	M
1012	F

Bitmap Index

Female	Male
1	0
1	0
0	1
1	0
1	0
0	1
1	0
0	1
1	0
0	1
0	1
1	0



# Bitmap Indexes – Many Distinct Values

Table Data

Customer ID	Age Range
1001	18-34
1002	18-34
1003	35-49
1004	35-49
1005	18-34
1006	62+
1007	50-61
1008	35-49
1009	35-49
1010	35-49
1011	35-49
1012	18-34



Bitmap Index

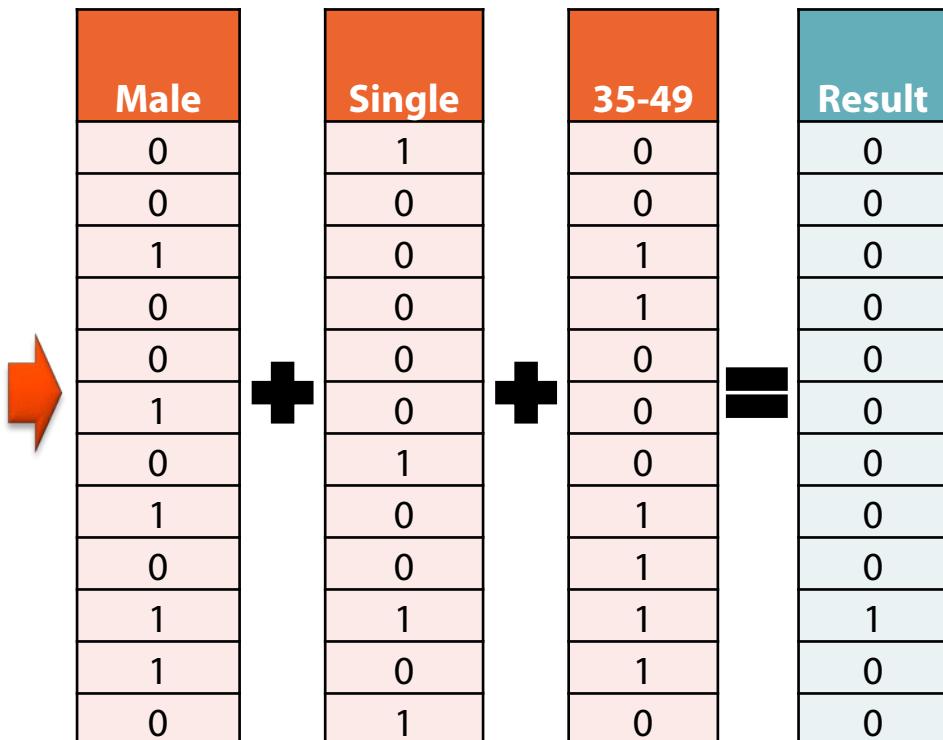
18-34	35-49	50-61	62+
1	0	0	0
1	0	0	0
0	1	0	0
0	1	0	0
1	0	0	0
0	0	0	1
0	0	1	0
0	1	0	0
0	1	0	0
0	1	0	0
1	0	0	0

# Using Bitmap Indexes

Table data

Customer ID	Gender	Married	Age Range
1001	F	N	18-34
1002	F	Y	35-49
1003	M	Y	35-49
1004	F	Y	35-49
1005	F	Y	18-34
1006	M	Y	62+
1007	F	N	50-61
1008	M	Y	35-49
1009	F	Y	35-49
1010	M	N	35-49
1011	M	Y	35-49
1012	F	N	18-34

Bit arrays from bitmap indexes



# Using Bitmap Indexes

Use for columns with low number of distinct values (few hundred max)

Use only in data warehouse or read only databases, not OLTP databases

Only available in  
Oracle Enterprise Edition

# Defining a Bitmap Index

```
CREATE BITMAP INDEX bx_customers_stats_gender  
    ON customer_statistics  
    (gender);
```

// More flexible to create two indexes like this

```
CREATE BITMAP INDEX bx_customer_stats_age_range  
    ON customer_statistics  
    (age_range);
```

// Allowed, not as flexible

```
CREATE BITMAP INDEX bx_customer_stats_gender_age  
    ON customer_statistics  
    (gender, age_range);
```

**Is there some advice I should know when  
defining indexes for my database?**

# Column Order Matters in an Index

## Index columns

last_name	first_name	state	city
-----------	------------	-------	------



Leading edge  
of the index

- ◆ SQL statement must use the first column to use index
- ◆ Use as many columns as possible from front of index
- ◆ Put columns most frequently used in WHERE clause at the front of the index

# Index Selectivity

## Moderate selectivity

last\_name

## Better selectivity

last\_name  
first\_name

## Superb selectivity

last\_name  
first\_name  
zip\_code



*More columns or more distinct columns  
help us identify the data more uniquely*

# Selectivity Example

## Gender column

---

- Only two distinct values
- Oracle will need to read half the index, then most likely the entire table
- More efficient to use a full table scan (don't bother to read the index)

## Email column

---

- Every row has a unique value
- B-tree structure very effective at efficiently locating unique or nearly unique values
- Data is found reading a minimum number of blocks

# Do Not Overindex Your Database

## Each index is a separate data structure in Oracle

---

- Every index takes up space on disk
- Oracle must maintain the index during DML operations against the table

## Look at indexes like an investment

---

- A good index repays its maintenance cost many times over by frequently speeding up other SQL statements
- A poor index is one you pay the cost for but never reap any dividends

# **Module Summary**

What is an index?

How do I create an index?

What is a function based index?

What is a bitmap index?

What are some indexing tips to follow?

# Privileges and Roles

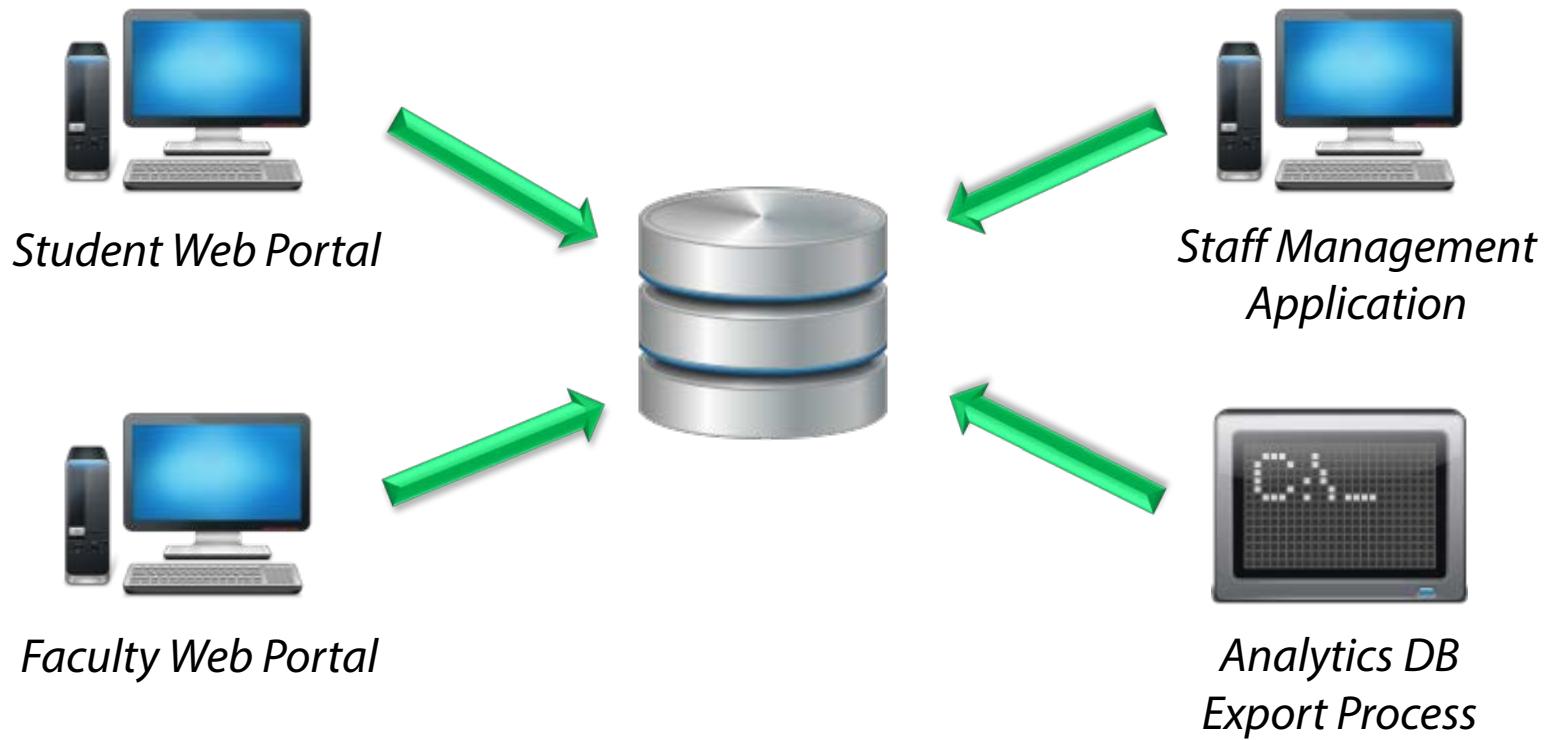
David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# What Applications Use Our Database?



*Should all of these applications connect  
as the user 'student'?*

# Sharing Users – Problems it Poses

## Full Access to All Objects

Application could read, write data it is not supposed to

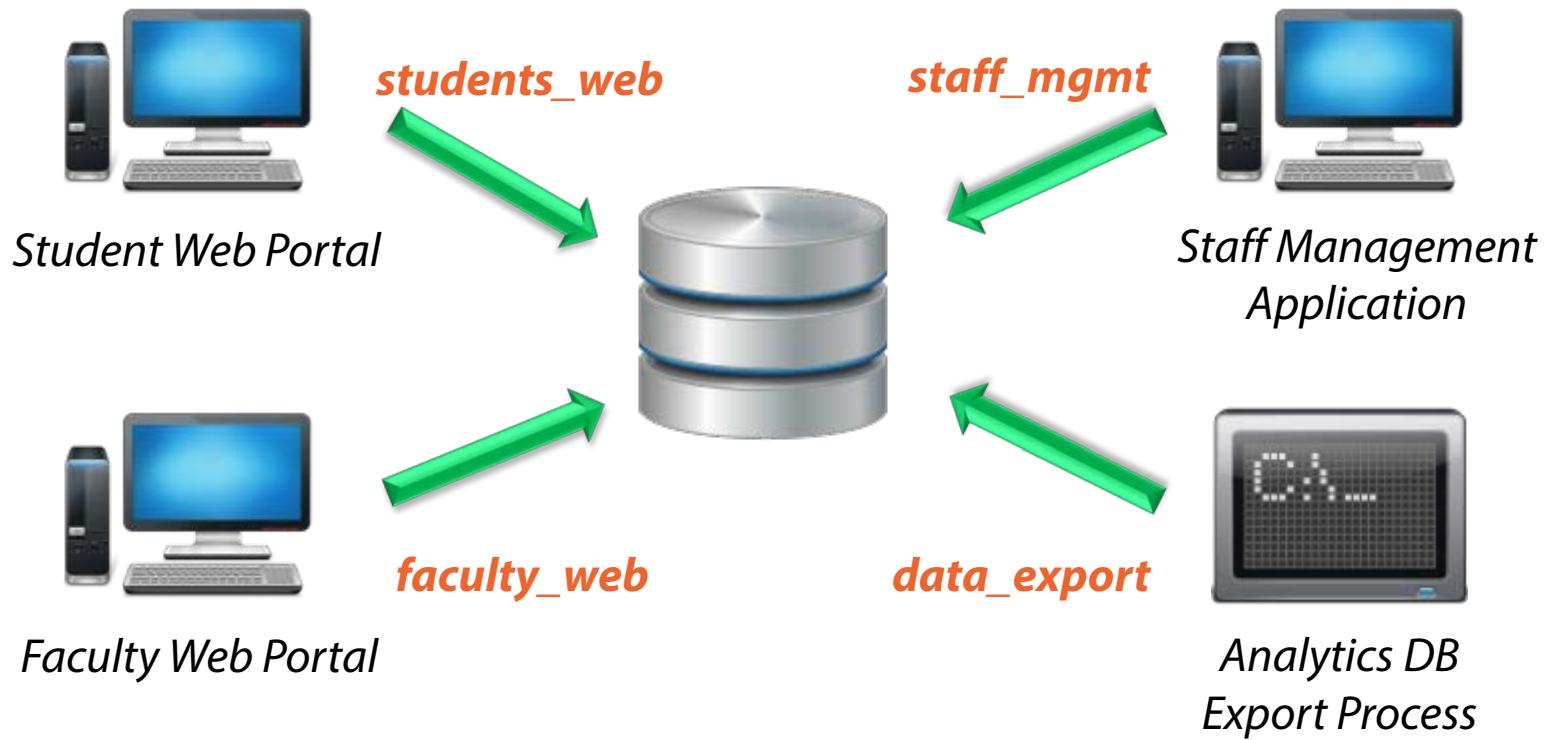
## Schema Changes

Should an application be adding, dropping tables on the fly?

## Who is Running SQL

Possible to determine, easier if separate user for each application

# What Applications Use Our Database?



# **Module Outline**

**GRANT and REVOKE**

**Object privileges**

**Column level permissions**

**Stored procedures and security**

**Database roles**

**Security principles**

# **GRANT and REVOKE Commands**

**GRANT**

Give a user privileges on an object

**REVOKE**

Remove privileges on an object

# **GRANT Command**

- **Basic syntax**

```
GRANT privileges ON object TO user;
```

- **Examples**

```
GRANT select ON courses TO students_web, faculty_web;
```

```
GRANT select, update ON students TO students_web;
```

# REVOKE Command

- **Basic syntax**

```
REVOKE privileges ON object FROM user;
```

- **Examples**

```
REVOKE update ON students FROM students_web, faculty_web;
```

```
REVOKE insert ON courses FROM faculty_web;
```

# Table Privileges

Privilege	Description
ALTER	Change the table definition using ALTER TABLE
DEBUG	Debug the PL/SQL in a trigger on the table
DELETE	Remove rows on the table via a DELETE statement
INDEX	Create an index on the table
INSERT	Add new rows to the table via an INSERT statement
READ	Allows queries from the table but not a SELECT ... FOR UPDATE statement (pessimistic locking)
REFERENCES	Allows a foreign key constraint to the table (can only be granted to a user, not a role)
SELECT	Allows reads from the table including SELECT ... FOR UPDATE statements
UPDATE	Modify rows in the table via an UPDATE statement

# View Privileges

Privilege	Description
DEBUG	Debug the PL/SQL in a trigger on the view
DELETE	Remove rows on the view via a DELETE statement
INSERT	Add new rows to the view via an INSERT statement
READ	Allows queries from the view but not a SELECT ... FOR UPDATE statement (pessimistic locking)
REFERENCES	Allows a foreign key constraint to the view (can only be granted to a user, not a role)
SELECT	Allows reads from the view including SELECT ... FOR UPDATE statements
UPDATE	Modify rows in the view via an UPDATE statement

# Column Level Privileges

```
GRANT select, update ON students TO students_web;
```



*User students\_web can **update any column**  
on the students table*

# Column Level Privileges

```
GRANT update (street_address, city, state, telephone, email)  
ON students TO students_web;
```



*User students\_web can only update  
columns **included in this list***

- Individual columns can be supplied for the insert, update and references privileges

# Using a View to Restrict Visible Columns

**Problem**

- A table contains columns I do not want another user to see

**Column Level  
Privileges**

- Not available for the select privileges

**Solution**

- Create a view with only the columns you want to expose
- Grant select access to the view

# Example – Employee Phone Directory

```
desc employees;
```

Name	Null	Type
EMPLOYEE_ID	NOT NULL	NUMBER(8),
FIRST_NAME	NOT NULL	VARCHAR2(30),
LAST_NAME	NOT NULL	VARCHAR2(30),
DATE_OF_BIRTH	NOT NULL	DATE,
TITLE	NOT NULL	VARCHAR2(30),
DEPARTMENT_ID	NOT NULL	NUMBER(4),
PHONE	NOT NULL	VARCHAR2(20),
EMAIL	NOT NULL	VARCHAR2(60),
SALARY	NOT NULL	NUMBER(10,2)

# Example – Employee Phone Directory

```
CREATE VIEW v_employee_contact_info AS
    SELECT
        employee_id, first_name, last_name,
        title, department_id, phone, email
    FROM employees;

GRANT select ON v_employee_contact_info TO web_user;
```

# **Stored Procedure/Function Privileges**

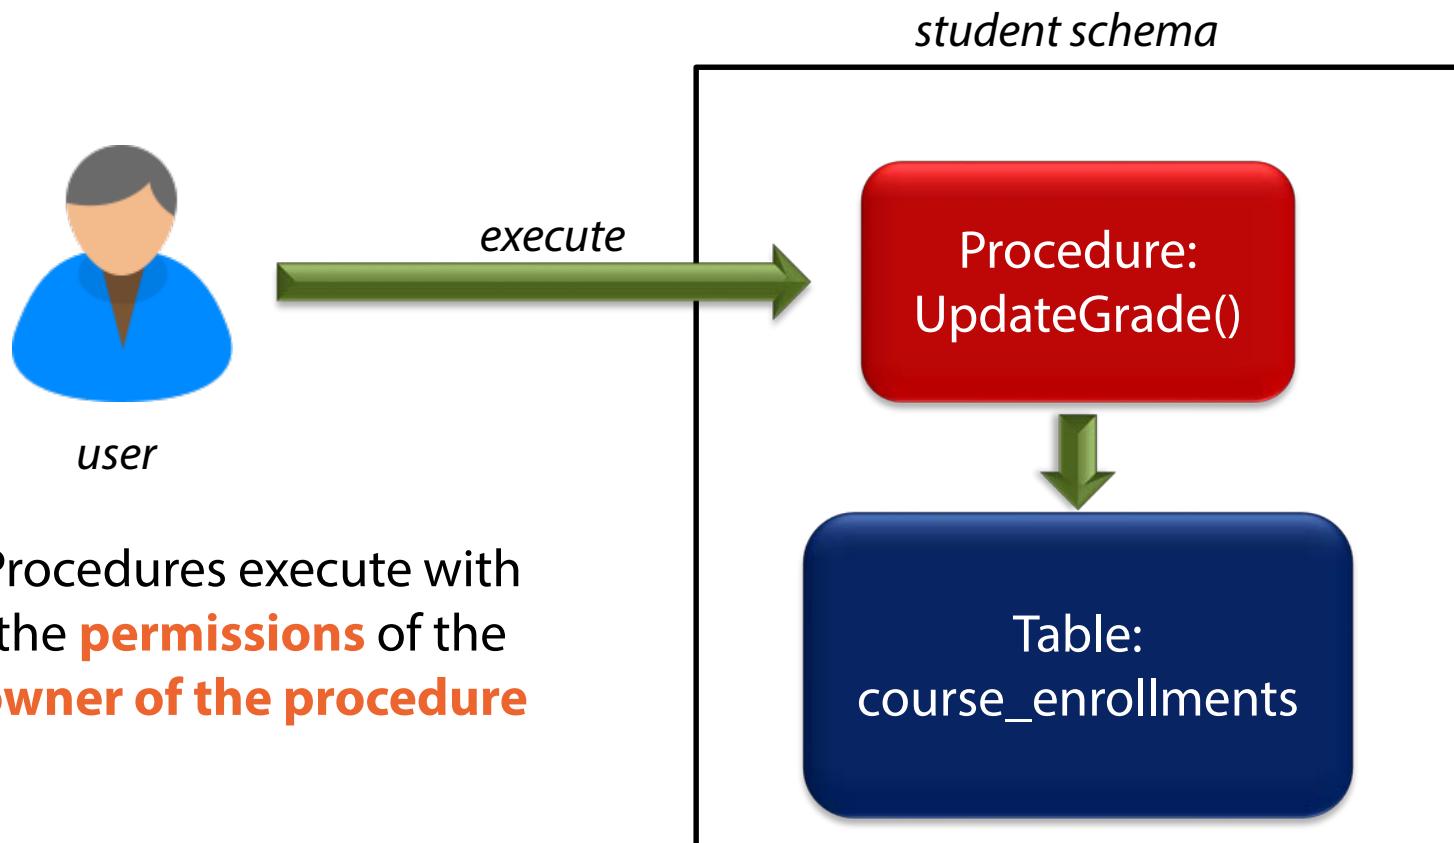
**EXECUTE**

Access to run the procedure or function

**DEBUG**

Access to debug through the procedure/function

# Procedure Execution Permissions



# **SQL in Applications**

## **Versus**

# **Stored Procedures**

Debated numerous times

No need to debate again here

Procedures can offer security benefits

# Direct Table Access and Security Concerns



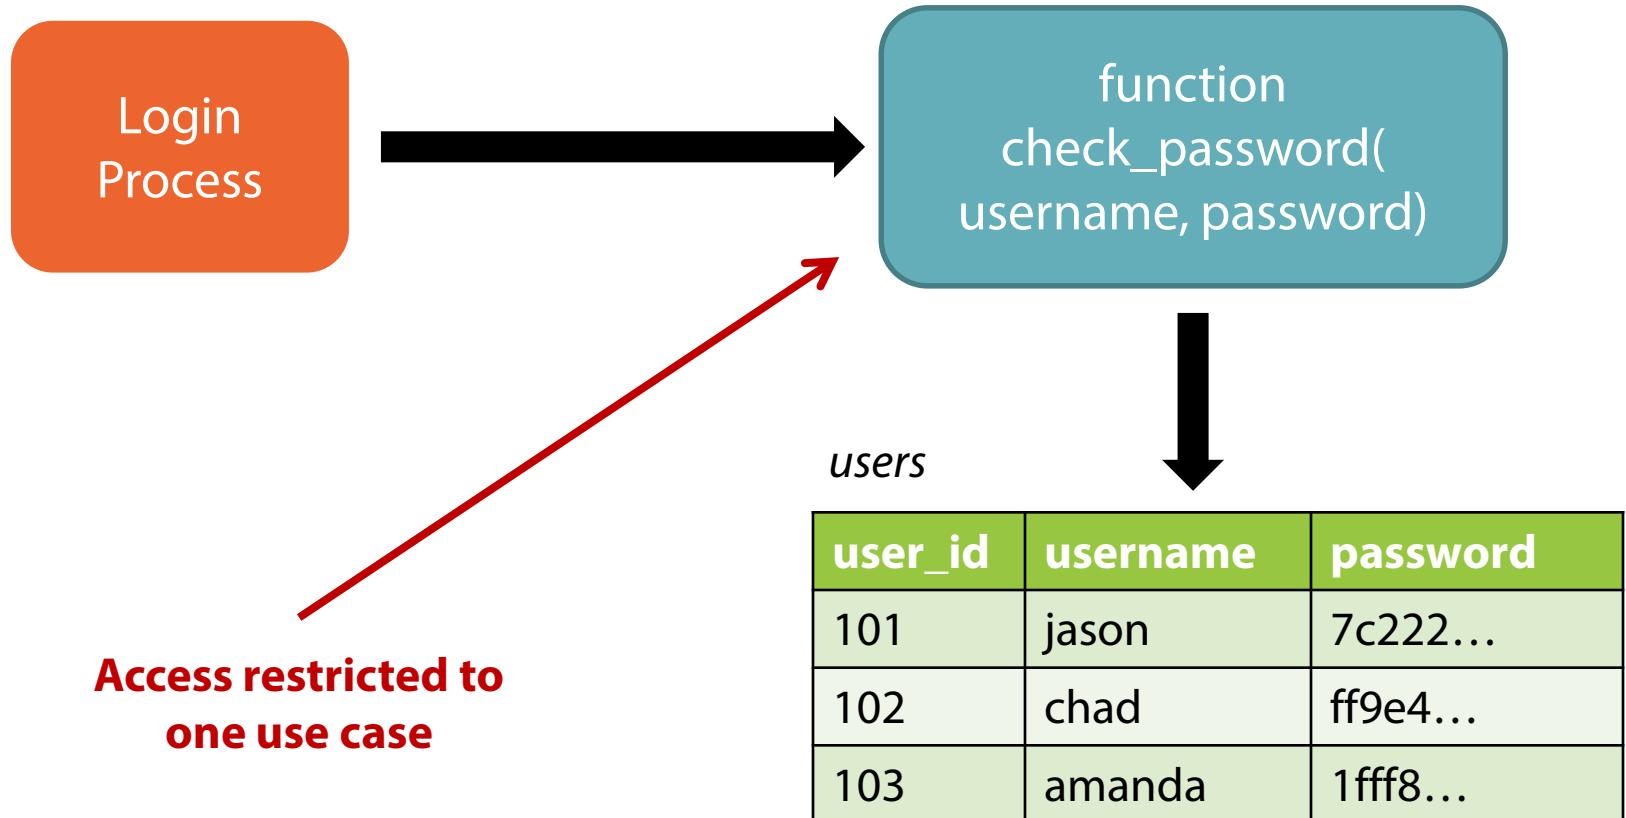
```
SELECT user_id  
FROM users  
WHERE username = ?  
AND password = ?
```



**Direct access to the table**

users		
user_id	username	password
101	jason	7c222...
102	chad	ff9e4...
103	amanda	1fff8...

# Stored Procedures and Security



# Stored Procedure Wrap Up

Grant **EXECUTE** privilege to allow users to run

Users only need EXECUTE privilege, **not** access to **underlying objects**

Can **restrict access** to a specific use case

Consider using when you have **sensitive data**

# Roles Introduction

```
GRANT select ON departments TO john;
```

```
GRANT select ON courses TO john;
```

```
GRANT select ON students TO john;
```

```
...
```

```
GRANT select ON departments TO erica;
```

```
GRANT select ON courses TO erica;
```

```
GRANT select ON students TO erica;
```

```
...
```

```
GRANT select ON departments TO robert;
```

```
GRANT select ON courses TO robert;
```

```
GRANT select ON students TO robert;
```

```
...
```

# Database Roles

Definition

Container to group a set of privileges together

Benefit

Simplify management of permissions

# Role Management

```
-- Create the role
-- Most likely will need to be done by a DBA
CREATE ROLE technology_staff;

-- Assign privileges to role
GRANT select ON departments TO technology_staff;
GRANT select ON courses TO technology_staff;
GRANT select ON students TO technology_staff;
...

-- Assign the role to users
-- Again, performed by a DBA
GRANT technology_staff TO john;
GRANT technology_staff TO erica;
GRANT technology_staff TO robert;
```

# Thoughts on Database Privileges

## Principle of Least Privilege

---

*Only grant access to the data that an application or user absolutely needs and no more*

## Defense in Depth

---

*Use multiple mechanisms to secure access to your data*

# Some Good Practices

## Separate application users from the schema owner

---

*Schema owner has full access to all objects*

*Schema owner can perform DDL statements*

## Use a separate Oracle user for each application

---

*Applications usually have different access requirements*

*Easier to tell what application a statement was executed from*

## Take time to assign appropriate access permissions

---

*Don't rely solely on application code to enforce security*

*Only give a user the access they need*

# Summary

**Multiple users**

**GRANT and REVOKE**

**Object privileges**

**Column level permissions**

**Stored procedures and security**

**Database roles**

**Security principles**

# Good Practices

David Berry

<http://buildingbettersoftware.blogspot.com/>



**pluralsight**   
hardcore dev and IT training

# Module Overview

-  Naming standards and conventions
-  Table and column comments
-  Single responsibility principle

# **Naming Standards and Conventions**

Establish and use standards for your organization

Consistent use makes your database easier to work with

Standards add an element of professionalism

# Naming Tables and Columns



Use a descriptive name of what data is contained in the table or column



Table names are commonly nouns



Be consistent in use of singular or plural forms of names



Column names should be specific about what attribute they represent

# Favor Full Words Over Abbreviations

Full words are **more descriptive**

**Abbreviations** are open to **misinterpretation**

Common abbreviations in **one culture** may be  
**confusing to others**

# User Underscores to Improve Readability

*Which set of names is easier to read?*

---

STREET\_ADDRESS

STREETADDRESS

DEPARTMENT\_NAME

DEPARTMENTNAME

REQUEST\_ID

REQUESTID

COURSE\_TITLE

COURSETITLE

# Name Your Table Constraints

Constraint type	Naming convention
Primary Key	PK_<table name>
Foreign Key	FK_<parent table name>_<parent column>
Check Constraint	CK_<table name>_<column name>

# Naming Indexes

**<prefix>\_<table name>\_<column name(s)>**



*Based on index type*

---

- IX      Non-unique index
- UX      Unique index
- FX      Function based index

***Goal: Look at an index and be able to tell its purpose***

## **Most Important Rule for Naming Objects**

---

*Be expressive*

## **Database Comments**

Needed when a concept is not self explanatory

Documentation works best when it can be stored with the structure it is documenting

# Adding Comments to Tables and Columns

```
-- Setting comments on a table  
COMMENT ON TABLE <<table name>> IS ‘<<comment test>>’;  
  
-- Setting a comment on a columns  
COMMENT ON COLUMN <<table name>>. <<column name>>  
IS ‘<<comment text>>’;
```

# Selecting Comments About Tables and Columns

```
SELECT * FROM user_tab_comments;

SELECT * FROM user_col_comments
  WHERE table_name = '<




```

## **Single Responsibility Principle**

---

*Every structure in your database  
should do exactly one thing*

**Tables should store a single type of entity**

# Do Some Columns Apply Only for Some Rows?

PERSON ID	FIRST NAME	LAST NAME	PERSON TYPE	EMAIL	ACADEMIC DEPT	TITLE	MAJOR	STATUS
1001	John	Randolph	S	jrandolph@...			CS	G
1002	Erica	Schmidt	S	eschmidt@...			EE	G
1003	Richard	Rodgers	S	rrodgers@...			Math	W
1004	Jennifer	Meehan	F	jmeehan@...	Math	Professor		
1005	Martha	McNeil	F	mmcneil@...	Physics	Assoc Professor		
1006	Eric	O'Reilly	S	eoreilly@...			EE	E
1007	William	Anderson	F	wanderson@...	Chemistry	Professor		
1008	Matthew	Taylor	S	mtaylor@...			ME	E

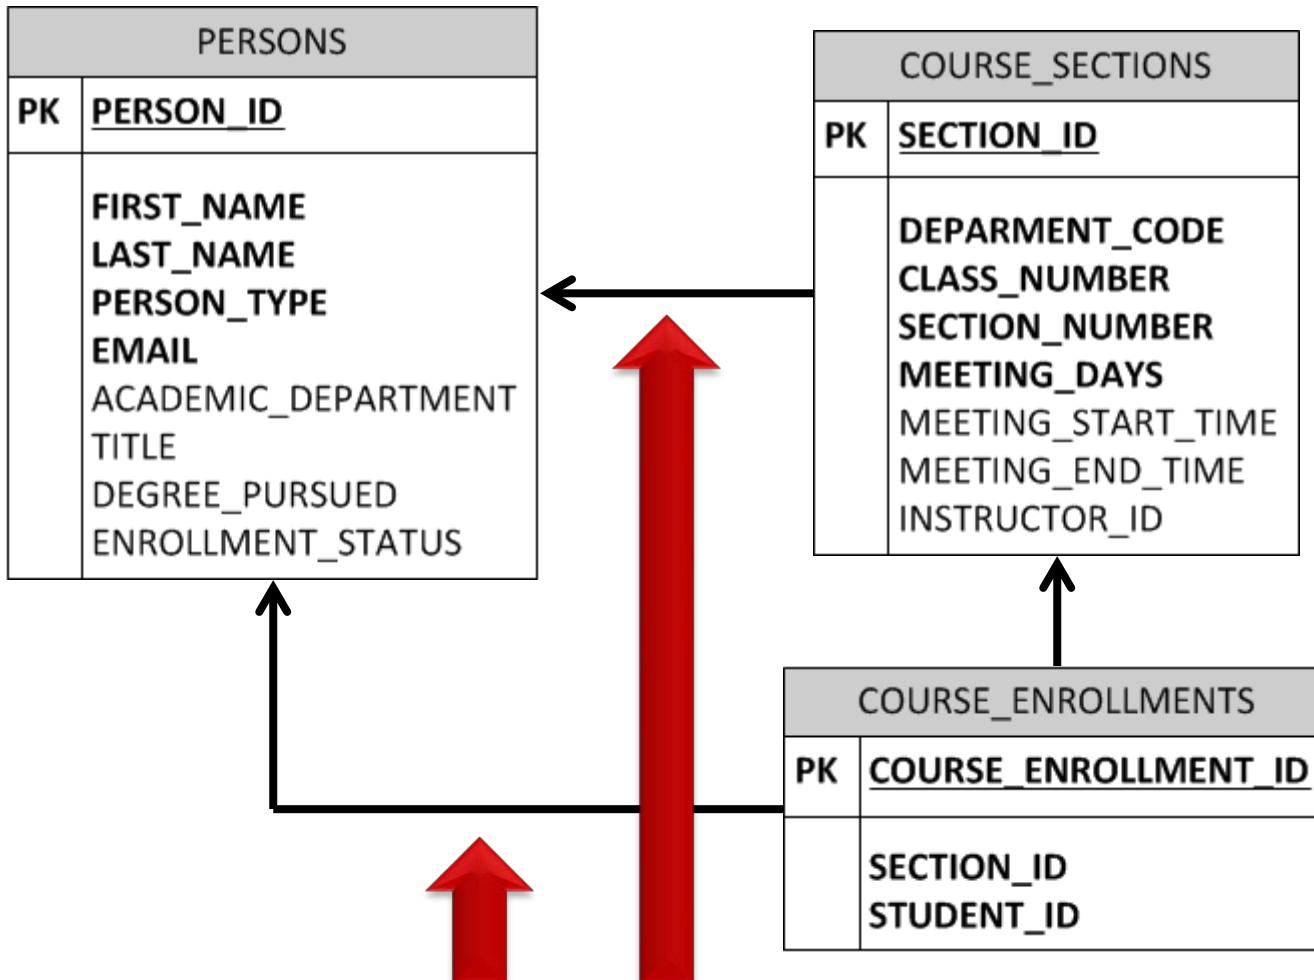


*Attributes apply to all rows*

*Only for faculty members*

*Only for students*

# Foreign Key Issues



**Problem: Cannot create foreign key to only faculty or only student records**

# Columns Should Have One Use

ORDER_ID	SHIP_STATUS	SHIP_DATE
1001	S	2014-07-15
1002	S	2014-07-17
1003	S	2014-07-18
1004	P	2014-07-25
1005	P	2014-07-27



Actual ship dates

Expected ship dates

# A Better Design

ORDER_ID	SHIP_STATUS	TARGET_SHIP_DATE	ACTUAL_SHIP_DATE
1001	S	2014-07-15	2014-07-15
1002	S	2014-07-18	2014-07-17
1003	S	2014-07-19	2014-07-18
1004	P	2014-07-25	
1005	P	2014-07-27	

## **Single Responsibility Principle Benefits**

Easier to understand the design

Easier to access data in tables

Easier to achieve good design

# **Summary**

Naming  
conventions

Database  
comments

Single  
responsibility  
principle